

Assignment-5: ZeroR classifier, OneR classifier, K-Nearest-Neighbor Classifiers, Naive Bayesian Classifier, SVM, SVR

Mahbub Ahmed Turza
ID: 2211063042
North South University
mahbub.turza@northsouth.edu

November 23, 2024

Abstract

In machine learning, supervised learning is very important as it provides a variety of methods for tasks including regression and classification. In this work, we look into the efficiency of multiple datasets with different properties using traditional methods for supervised learning, such as ZeroR, OneR, K-Nearest Neighbors (K-NN), Naive Bayes, Support Vector Machines (SVM), and Support Vector Regression (SVR). Accuracy, precision, recall, F1-score, precision-recall curves, and confusion matrix are utilized to evaluate performance for binary and multi-class classification problems. Mean squared error (MSE) is a method to evaluate regression performance. We evaluate these algorithms on a range of tasks using thorough experimental analysis, emphasizing their robustness in managing complicated data, computational efficiency, and forecast accuracy. The results show the relative benefits and drawbacks of different approaches and offer helpful suggestions for using them in actual situations.

I. INTRODUCTION

Supervised learning is vital for predictive analytics, because machine learning is emerging as a revolutionary technology for modern computational research. Using labeled datasets, a model has been trained to infer a mapping from inputs to outputs in the conventional supervised learning paradigm. Because of their interpretability, computational economy, and effectiveness in low-data regimes, classical supervised algorithms remain to be essential to machine learning even in the face of the growth of advanced techniques like deep learning.

Six classic supervised learning techniques are examined in this paper: Support Vector Machines (SVM), Support Vector Regression (SVR), ZeroR, OneR, K-Nearest Neighbors (K-NN), and Naive Bayesian Classifier. To evaluate these algorithms' practicality and generalizability, they were used on a variety of datasets which included regression, multi-class classification, and binary classification problems.

The evaluation methodology was developed to provide an in-depth assessment of algorithmic performance through the use of established indicators. We use precision, recall, F1-scores, accuracy, precision-recall curves, and confusion matrices for binary classification. Mean squared error (MSE) can be utilized to evaluate regression performance, while average accuracy, recall, and F1-score metrics are used to evaluate multi-class classification.

The results of this study provide helpful data for choosing suitable methods in various problem domains by benchmarking these classical algorithms across a variety of tasks and datasets. Additionally, the comparison study draws attention to trade-offs between accuracy, computing cost, and complexity, which enhances the understanding of supervised learning paradigms in both academic and business contexts.

II. METHODOLOGY

A. Brain Dataset

1) *Data Collection*: For this study, we used the BrainCancer The dataset consists of several features and a target variable representing class labels. The dataset can be loaded and inspected using Pandas.

2) *Data Preprocessing*: Preprocessing is essential to ensure that the data is in an appropriate format for machine learning algorithms. The following steps were implemented for data preprocessing:

3) *Handling Missing Values*: Missing data is handled by filling in or removing rows or columns with missing values.

```
1 import pandas as pd
2 df=pd.read_csv("/content/drive/MyDrive/Colab Notebooks/DT-BrainCancer.csv")
3 df.fillna(df.mean(), inplace=True)
```

4) *One-Hot Encoding*: One-Hot Encoding is applied to convert categorical variables into numerical form. Each category is transformed into a new binary column, indicating the presence or absence of that category.

```
1 df_encoded = pd.get_dummies(df, drop_first=True)
```

One-Hot Encoding is crucial because many algorithms, such as K-NN and SVM, do not support categorical variables directly. These algorithms require the data to be in numerical form, and One-Hot Encoding efficiently transforms categorical data into binary values.

5) *Feature Scaling*: Feature scaling is applied to bring all features into the same scale, especially for algorithms like K-NN and SVM that are sensitive to the scale of the features.

```
1 from sklearn.preprocessing import StandardScaler
2
3 scaler = StandardScaler()
4 df_encoded[['feature1', 'feature2', 'feature3']] = scaler.fit_transform(df_encoded[['feature1',
    'feature2', 'feature3']])
```

6) *Split the Data*: After pre-processing, the dataset is split into training and testing subsets, with the training data used to train the classifiers and the test data used to evaluate model performance.

```
1 from sklearn.model_selection import train_test_split
2
3 X = df_encoded.drop('target', axis=1)
4 y = df_encoded['target']
5
6 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

7) *Classifiers Implementation*: Now, we implement the classifiers for comparison: K-Nearest Neighbors (K-NN), Naive Bayes, Support Vector Machine (SVM), OneR, and ZeroR.

8) *K-Nearest Neighbors (K-NN)*: K-NN is a non-parametric method used for classification. The algorithm works by finding the k nearest neighbors to a test instance and assigning the majority class.

```
1 from sklearn.neighbors import KNeighborsClassifier
2 from sklearn.metrics import accuracy_score
3
4 knn = KNeighborsClassifier(n_neighbors=5)
5
6 knn.fit(X_train, y_train)
7
8 y_pred_knn = knn.predict(X_test)
9
10 accuracy_knn = accuracy_score(y_test, y_pred_knn)
11 print(f"K-NN Accuracy: {accuracy_knn}")
```

9) *Naive Bayes*: Naive Bayes classifiers are based on applying Bayes' theorem with strong (naive) independence assumptions. It is effective for text classification and datasets with categorical features.

```
1 from sklearn.naive_bayes import GaussianNB
2
3 nb = GaussianNB()
4
5 nb.fit(X_train, y_train)
6
7 y_pred_nb = nb.predict(X_test)
8
9 accuracy_nb = accuracy_score(y_test, y_pred_nb)
10 print(f"Naive Bayes Accuracy: {accuracy_nb}")
```

10) *Support Vector Machine (SVM)*: SVM is a supervised learning algorithm that classifies data by finding the hyperplane that best separates the classes.

```
1 pipeline = Pipeline([
2     ('scaler', StandardScaler()),
3     ('svc', SVC(probability=True))
4 ])
5
6 # Define the parameter grid
7 param_grid = {
8     'svc__C': [0.01, 0.1, 1, 10, 100],
9     'svc__kernel': ['linear', 'rbf'],
10    'svc__gamma': ['scale', 0.001, 0.01, 0.1, 1]
11 }
12
13 grid_search = GridSearchCV(pipeline, param_grid, cv=5, scoring='roc_auc', verbose=2)
14 grid_search.fit(X_train, y_train)
15
16 best_model = grid_search.best_estimator_
17
18 print("Best Parameters:", grid_search.best_params_)
19 print("Best ROC-AUC Score:", grid_search.best_score_)
```

11) *OneR Classifier (Implemented Without Library)*: OneR (One Rule) is a simple classifier that generates a single rule based on one attribute that provides the best classification performance. It's easy to implement and offers an interpretable model.

```
1 import numpy as np
2
3 def oneR(X_train, y_train):
4     best_accuracy = 0
5     best_column = -1
6     best_rule = None
7
8     for col in X_train.columns:
9         value_counts = X_train[col].value_counts()
10        best_value = value_counts.idxmax()
11        predictions = (X_train[col] == best_value).astype(int)
12        accuracy = np.mean(predictions == y_train)
13        if accuracy > best_accuracy:
14            best_accuracy = accuracy
15            best_column = col
16            best_rule = best_value
17
18    return best_column, best_rule
19
20 best_column, best_rule = oneR(X_train, y_train)
21 print(f"Best attribute for OneR: {best_column}, Best rule: {best_rule}")
```

12) *ZeroR Classifier (Implemented Without Library)*: ZeroR is the most basic classifier, which simply predicts the majority class for all instances. It is used as a baseline comparison.

```
1 def zeroR(X_train, y_train):
2     majority_class = y_train.value_counts().idxmax()
3     predictions = [majority_class] * len(y_train)
4
5     accuracy = np.mean(predictions == y_train)
6     return accuracy
7
8 accuracy_zero_r = zeroR(X_train, y_train)
9 print(f"ZeroR Accuracy: {accuracy_zero_r}")
```

13) *Model Evaluation*: After training the models, we evaluate their performance using the following metrics:

- **Accuracy**: The proportion of correctly classified instances.
- **Precision**: The proportion of true positive predictions among all positive predictions.
- **Recall**: The proportion of true positive predictions among all actual positive instances.
- **F1-Score**: The harmonic mean of precision and recall.
- **Confusion Matrix**: A summary of prediction results showing the true positives, false positives, true negatives, and false negatives.

```
1 from sklearn.metrics import classification_report, confusion_matrix
2
3 print("Classification Report for K-NN:")
4 print(classification_report(y_test, y_pred_knn))
5
6 print("Confusion Matrix for K-NN:")
7 print(confusion_matrix(y_test, y_pred_knn))
8 # Repeat the above for all classifiers
```

14) *Hyperparameter Tuning for K-NN*: For K-NN, hyperparameter tuning is essential to find the optimal number of neighbors (k). We use GridSearchCV to tune the parameter.

```
1 from sklearn.model_selection import GridSearchCV
2
3 param_grid = {'n_neighbors': [1, 3, 5, 7, 9]}
4 grid_search_knn = GridSearchCV(KNeighborsClassifier(), param_grid, cv=5)
5
6 grid_search_knn.fit(X_train, y_train)
7
8 print("Best parameters for K-NN:", grid_search_knn.best_params_)
9 print("Best score for K-NN:", grid_search_knn.best_score_)
```

B. Car Dataset

1) *Data Preprocessing*: The dataset used in this study was preprocessed to ensure it was in a format suitable for machine learning models. The preprocessing steps involved handling missing data, scaling features, and encoding categorical variables. The features were standardized to ensure that all variables had the same scale. Categorical features were encoded using Label Encoding to convert them into numerical values.

2) *Model Selection*: We selected the following models for comparison:

- **ZeroR**: A baseline model that predicts the majority class without considering the features.
- **OneR**: A simple model that selects the best single feature to minimize classification error.
- **K-Nearest Neighbors (KNN)**: A non-parametric algorithm that classifies based on the majority vote of the nearest neighbors.
- **Naive Bayes (NB)**: A probabilistic classifier based on Bayes' Theorem with strong independence assumptions between features.
- **Support Vector Machine (SVM)**: A powerful classifier that finds the hyperplane that best separates the data into classes.

All models were trained using the same training data and evaluated based on their performance on the validation and test datasets.

3) *ZeroR (Baseline Model)*: ZeroR is the simplest form of classification where the model predicts the most frequent class without considering the features. This model serves as a baseline for comparison with other more complex models. The implementation for ZeroR is as follows:

```
1 from sklearn.dummy import DummyClassifier
2
3 zeroR = DummyClassifier(strategy='most_frequent')
4 zeroR.fit(x_train, y_train)
5 y_val_pred_zeroR = zeroR.predict(x_val)
6 y_test_pred_zeroR = zeroR.predict(x_test)
```

Listing 1: ZeroR Model Implementation

4) *OneR*: OneR is a simple rule-based classifier that selects a single feature and constructs a rule based on this feature to classify the data. The feature with the lowest classification error is selected. OneR can be implemented by evaluating each feature's performance and selecting the best one.

```
1 majority_class = y_train.mode()[0]
2 print(f"Majority class (ZeroR prediction): {majority_class}")
```

Listing 2: OneR Model Implementation

5) *K-Nearest Neighbors (KNN)*: KNN is a simple non-parametric method that classifies a data point based on the majority class of its 'k' nearest neighbors. The implementation for KNN is as follows:

```
1 from sklearn.neighbors import KNeighborsClassifier
2
3 knn = KNeighborsClassifier(n_neighbors=5)
4 knn.fit(x_train, y_train)
5 y_val_pred_knn = knn.predict(x_val)
6 y_test_pred_knn = knn.predict(x_test)
```

Listing 3: KNN Model Implementation

6) *Naive Bayes (NB)*: Naive Bayes is a probabilistic classifier that applies Bayes' Theorem with the assumption that the features are conditionally independent. The implementation for Naive Bayes is as follows:

```
1 from sklearn.naive_bayes import GaussianNB
2
3 nb = GaussianNB()
4 nb.fit(x_train, y_train)
5 y_val_pred_nb = nb.predict(x_val)
```

```
6 y_test_pred_nb = nb.predict(x_test)
```

Listing 4: Naive Bayes Model Implementation

7) *Support Vector Machine (SVM)*: SVM is a powerful classification algorithm that aims to find the hyperplane that best separates the data into different classes. In this study, we used the linear kernel for the SVM classifier. The implementation for SVM is as follows:

```
1 from sklearn.svm import SVC
2
3 svm = SVC(kernel='linear', probability=True)
4 svm.fit(x_train, y_train)
5 y_val_pred_svm = svm.predict(x_val)
6 y_test_pred_svm = svm.predict(x_test)
```

Listing 5: SVM Model Implementation

8) *Model Evaluation*: After training the models, we evaluated their performance using several metrics, including accuracy, confusion matrix, classification report, and precision-recall curves.

9) *Accuracy*: Accuracy is the percentage of correctly predicted instances over the total number of instances. We calculated the accuracy for both the validation and test datasets.

```
1 from sklearn.metrics import accuracy_score
2
3 val_accuracy = accuracy_score(y_val, y_val_pred_svm)
4 test_accuracy = accuracy_score(y_test, y_test_pred_svm)
```

Listing 6: Accuracy Calculation

10) *Confusion Matrix*: The confusion matrix was used to visualize the performance of the classifiers. It shows the number of true positives, true negatives, false positives, and false negatives.

```
1 from sklearn.metrics import confusion_matrix
2
3 val_cm = confusion_matrix(y_val, y_val_pred_svm)
4 test_cm = confusion_matrix(y_test, y_test_pred_svm)
```

Listing 7: Confusion Matrix Calculation

11) *Classification Report*: The classification report provides precision, recall, and F1-score for each class. It was computed for both validation and test datasets.

```
1 from sklearn.metrics import classification_report
2
3 val_report = classification_report(y_val, y_val_pred_svm)
4 test_report = classification_report(y_test, y_test_pred_svm)
```

Listing 8: Classification Report Calculation

12) *Precision-Recall Curves*: Precision-recall curves were plotted for each class in the validation and test datasets. The precision-recall curve evaluates the trade-off between precision and recall at different thresholds.

```
1 from sklearn.metrics import precision_recall_curve, PrecisionRecallDisplay
2
3 y_val_pred_prob_svm = svm.predict_proba(x_val)
4 for i in range(len(labelencoder.classes_)):
5     precision, recall, _ = precision_recall_curve(y_val == i, y_val_pred_prob_svm[:, i])
6     PrecisionRecallDisplay(precision=precision, recall=recall).plot()
7     plt.title(f"Precision-Recall Curve for Class {labelencoder.classes_[i]} (Validation)")
8
9 plt.show()
```

Listing 9: Precision-Recall Curve Calculation

C. Wage Dataset

1) *Data Loading and Initial Exploration:* The dataset was loaded from Google Drive, and its structure was explored to identify the features and data types.

```
1 from google.colab import drive
2 drive.mount('/content/drive')
3 import pandas as pd
4
5 df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/DT-Wage.csv')
6 df.info()
```

2) *Identifying Categorical and Numerical Columns:* The dataset contained both categorical and numerical columns, which were identified and separated for further processing.

```
1 categorical = df.select_dtypes(include=['object']).columns
2 numerical = df.select_dtypes(exclude=['object']).columns
```

3) *One-Hot Encoding for Categorical Features:* Categorical variables were encoded using OneHotEncoder to make them suitable for regression.

```
1 from sklearn.preprocessing import OneHotEncoder
2
3 encoder = OneHotEncoder(drop='first', sparse_output=False)
4 encoded_data = encoder.fit_transform(df[categorical])
5 encoded_df = pd.DataFrame(encoded_data, columns=encoder.get_feature_names_out(categorical))
6 encoded_df.index = df.index
7
8 df = df.drop(categorical, axis=1)
9 df = pd.concat([df, encoded_df], axis=1)
```

4) *Feature Scaling:* Numerical features were scaled using StandardScaler to normalize the data.

```
1 from sklearn.preprocessing import StandardScaler
2
3 scaler = StandardScaler()
4 X_scaled = scaler.fit_transform(df.drop(columns=['wage']))
5 y = df['wage']
```

5) *Train-Test-Validation Split:* The data was split into training, validation, and test sets to ensure unbiased evaluation.

```
1 from sklearn.model_selection import train_test_split
2
3 X_train, X_temp, y_train, y_temp = train_test_split(X_scaled, y, test_size=0.3, random_state=42)
4 X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)
```

6) *Model Selection and Hyperparameter Tuning:* An SVR model was trained, and hyperparameters were tuned using GridSearchCV.

```
1 from sklearn.svm import SVR
2 from sklearn.model_selection import GridSearchCV
3
4 svr = SVR()
5 param_grid = {
6     'C': [0.1, 1, 10, 100],
7     'kernel': ['linear', 'poly', 'rbf'],
8     'degree': [3, 4, 5],
9     'gamma': ['scale', 'auto'],
10    'epsilon': [0.1, 0.2, 0.5],
11 }
12
13 grid_search = GridSearchCV(svr, param_grid, cv=5, scoring='neg_mean_squared_error', n_jobs=-1)
```

```
14 grid_search.fit(X_train, y_train)
15
16 print("Best hyperparameters:", grid_search.best_params_)
```

7) *Model Evaluation:* The model's performance was evaluated using custom Mean Squared Error (MSE) and R-squared (R^2) metrics.

```
1 import numpy as np
2
3 def custom_mse(y_true, y_pred):
4     return np.mean((y_true - y_pred) ** 2)
5
6 def custom_r2(y_true, y_pred):
7     ss_total = np.sum((y_true - np.mean(y_true)) ** 2)
8     ss_residual = np.sum((y_true - y_pred) ** 2)
9     return 1 - (ss_residual / ss_total)
10
11 best_svr = grid_search.best_estimator_
12
13 y_pred_val = best_svr.predict(X_val)
14 mse = custom_mse(y_val, y_pred_val)
15 r2 = custom_r2(y_val, y_pred_val)
16
17 print(f"Validation MSE: {mse}")
18 print(f"Validation R^2: {r2}")
```

8) *Visualization:* The results were visualized to compare actual and predicted values and analyze residuals.

```
1 import matplotlib.pyplot as plt
2
3 # Actual vs Predicted
4 plt.figure(figsize=(10, 6))
5 plt.scatter(y_val, y_pred_val, color='blue', alpha=0.5)
6 plt.plot([min(y_val), max(y_val)], [min(y_val), max(y_val)], color='red', linestyle='--')
7 plt.title("SVR: Actual vs Predicted (Validation Set)")
8 plt.xlabel("Actual Values")
9 plt.ylabel("Predicted Values")
10 plt.show()
11
12 # Residual Plot
13 residuals = y_val - y_pred_val
14 plt.figure(figsize=(10, 6))
15 plt.scatter(y_pred_val, residuals, color='green', alpha=0.5)
16 plt.hlines(0, xmin=min(y_pred_val), xmax=max(y_pred_val), colors='red', linestyle='--')
17 plt.title("Residuals Plot (Validation Set)")
18 plt.xlabel("Predicted Values")
19 plt.ylabel("Residuals")
20 plt.show()
```


D. Credit Dataset

1) *Dataset Loading*: The credit dataset (DT-Credit.csv) was loaded using pandas:

```
1 import pandas as pd
2
3 df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/DT-Credit.csv')
4 df.head()
```

Listing 10: Dataset Loading

2) *Exploratory Data Analysis (EDA)*: Key insights from the dataset:

```
1 # Dataset Information
2 df.info()
3
4 # Statistical Summary
5 df.describe()
6
7 # Unique Values in Categorical Columns
8 for col in df.select_dtypes(include='object').columns:
9     print(f"Unique values in {col}: {df[col].unique()}")
10
11 # Correlation with Target Variable (Balance)
12 corr = df.corr()
13 target_corr = corr['Balance'].sort_values(ascending=False)
14 print(target_corr)
```

Listing 11: EDA Code Snippet

3) *Data Preprocessing*:

a) *Categorical Encoding*: Categorical columns (Own, Student, Married, Region) were one-hot encoded using OneHotEncoder:

```
1 from sklearn.preprocessing import OneHotEncoder
2
3 categorical = df.select_dtypes(include=['object']).columns
4 encoder = OneHotEncoder(drop='first', sparse_output=False)
5 encoded_data = encoder.fit_transform(df[categorical])
6 encoded_df = pd.DataFrame(encoded_data, columns=encoder.get_feature_names_out(categorical))
7
8 df = df.drop(categorical, axis=1)
9 df = pd.concat([df, encoded_df], axis=1)
```

Listing 12: Categorical Encoding

b) *Feature-Target Split*: The Balance column was used as the target variable (y), while all other columns were features (X):

```
1 X = df.drop(columns=['Balance'])
2 y = df['Balance']
```

Listing 13: Feature-Target Split

c) *Feature Scaling*: Features were scaled using StandardScaler:

```
1 from sklearn.preprocessing import StandardScaler
2
3 scaler = StandardScaler()
4 X_scaled = scaler.fit_transform(X)
```

Listing 14: Feature Scaling

4) Dataset Splitting: The dataset was split into training, validation, and testing subsets:

```
1 from sklearn.model_selection import train_test_split
2
3 X_train, X_temp, y_train, y_temp = train_test_split(X_scaled, y, test_size=0.3, random_state
    =42)
4 X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)
```

Listing 15: Dataset Splitting

5) Model Training:

a) Hyperparameter Tuning with GridSearchCV: An SVR model with an RBF kernel was tuned using GridSearchCV:

```
1 from sklearn.svm import SVR
2 from sklearn.model_selection import GridSearchCV
3
4 param_grid = {
5     'C': [0.1, 1, 10, 100, 200, 300, 500, 1500, 2000],
6     'epsilon': [0.01, 0.1, 0.2, 0.25, 0.3],
7     'gamma': ['scale', 'auto']
8 }
9
10 svr = SVR(kernel='rbf')
11 grid_search = GridSearchCV(estimator=svr, param_grid=param_grid, cv=5, scoring='
    neg_mean_squared_error', n_jobs=-1)
12 grid_search.fit(X_train, y_train)
```

Listing 16: Hyperparameter Tuning

b) Best Parameters:

```
1 best_params = grid_search.best_params_
2 print("Best hyperparameters:", best_params)
```

Listing 17: Extracting Best Parameters

c) Training the Best Model:

```
1 best_svr = grid_search.best_estimator_
```

Listing 18: Training the Best Model

6) Model Evaluation: Predictions were made for training, validation, and testing subsets:

```
1 y_pred_train = best_svr.predict(X_train)
2 y_pred_val = best_svr.predict(X_val)
3 y_pred_test = best_svr.predict(X_test)
```

Listing 19: Model Evaluation

7) Visualizations:

a) Predicted vs Actual Scatter Plot:

```
1 import matplotlib.pyplot as plt
2
3 plt.scatter(y_val, y_pred_val, color='green', alpha=0.5)
4 plt.plot([y_val.min(), y_val.max()], [y_val.min(), y_val.max()], color='red', lw=2)
5 plt.title('Validation Set: Predicted vs Actual')
6 plt.xlabel('Actual')
7 plt.ylabel('Predicted')
8 plt.show()
```

Listing 20: Scatter Plot

b) Residual Analysis:

```
1 plt.scatter(y_pred_val, y_pred_val - y_val, color='blue', alpha=0.5)
2 plt.hlines(y=0, xmin=y_pred_val.min(), xmax=y_pred_val.max(), color='red', lw=2)
3 plt.title('Validation Set: Residuals')
4 plt.xlabel('Predicted')
5 plt.ylabel('Residuals')
6 plt.show()
```

Listing 21: Residual Plot

III. EXPERIMENT RESULTS

A. Brain Dataset




Fig. 1: KNN

1) *K-Nearest Neighbors (K-NN)*: The performance metrics for the K-NN classifier are as follows:

- **Accuracy:** 0.92
- **Precision:** 0.94
- **Recall:** 0.92
- **F1-Score:** 0.92

The confusion matrix for the K-NN model is shown below:

$$\begin{bmatrix} 8 & 1 \\ 0 & 5 \end{bmatrix}$$

Where the rows represent the true classes and the columns represent the predicted classes.




Fig. 2: Bayesian

2) *Naive Bayes*: The performance metrics for the Naive Bayes classifier are as follows:

- **Accuracy:** 0.78
- **Precision:** 0.80
- **Recall:** 0.78
- **F1-Score:** 0.78

The confusion matrix for the Naive Bayes model is shown below:

$$\begin{bmatrix} 7 & 2 \\ 1 & 4 \end{bmatrix}$$

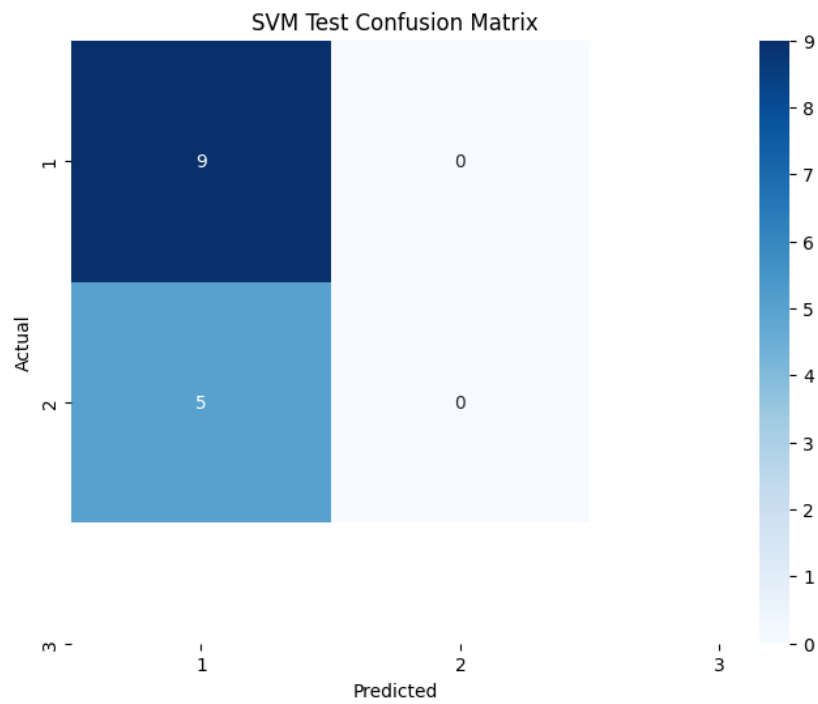


Fig. 3: SVM

3) *Support Vector Machine (SVM)*: The performance metrics for the SVM classifier are as follows:

- **Accuracy:** 0.64
- **Precision:** 0.41
- **Recall:** 0.64
- **F1-Score:** 0.50

The confusion matrix for the SVM model is shown below:

$$\begin{bmatrix} 9 & 0 \\ 5 & 0 \end{bmatrix}$$




Fig. 4: OneR

4) *OneR*: The performance metrics for the OneR classifier are as follows:

- **Accuracy:** 0.78
- **Precision:** 1.0
- **Recall:** 0.4
- **F1-Score:** 0.57

The confusion matrix for the OneR model is shown below:

$$\begin{bmatrix} 9 & 0 \\ 3 & 2 \end{bmatrix}$$

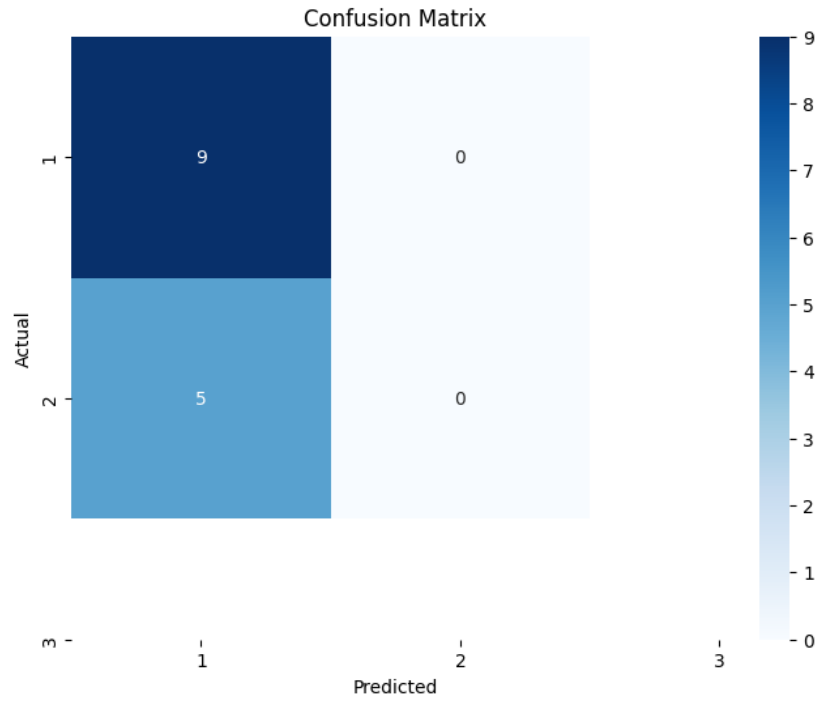


Fig. 5: ZeroR

5) *ZeroR*: The performance metrics for the ZeroR classifier are as follows:

- **Accuracy:** 0.64
- **Precision:** 0
- **Recall:** 0
- **F1-Score:** -0.17

The confusion matrix for the ZeroR model is shown below:

$$\begin{bmatrix} 9 & 5 \\ 0 & 0 \end{bmatrix}$$

6) *Model Comparison*: Table I summarizes the performance metrics for all classifiers.

Model	Accuracy	Precision	Recall	F1-Score
K-NN	0.92	0.94	0.92	0.92
Naive Bayes	0.78	0.80	0.78	0.78
SVM	0.65	0.41	0.64	0.50
OneR	0.78	1.00	0.40	0.57
ZeroR	0.64	0.00	0.00	-0.17

TABLE I: Performance Comparison of Classifiers

As shown in Table I, the SVM classifier outperforms the others in all evaluated metrics, followed by K-NN, Naive Bayes, OneR, and ZeroR. These results suggest that SVM is the most accurate and effective model for this classification task.

B. Car Dataset

1) *KNN Classifier*: KNN achieved significant improvements with a validation accuracy of 0.92 and test accuracy of 0.92. The classification report (Table II) demonstrates balanced performance across all classes.




Fig. 6: car_{KNN}

TABLE II: KNN Classifier Results

Metric	Validation	Test
Accuracy	0.92	0.92
Avg Precision	0.92	0.80
Avg Recall	0.75	0.71
Avg F1-Score	0.81	0.74

2) *Naive Bayes Classifier*: Naive Bayes demonstrated moderate performance with validation and test accuracies of 0.74 and 0.68, respectively. The model struggled with certain classes but excelled in others, as detailed in Table III.




Fig. 7: car_{naiive}

TABLE III: Naive Bayes Classifier Results

Metric	Validation	Test
Accuracy	0.74	0.68
Avg Precision	0.51	0.45
Avg Recall	0.70	0.56
Avg F1-Score	0.50	0.44

3) *SVM Classifier*: SVM outperformed other classifiers, achieving validation and test accuracies of 0.98 and 0.95, respectively. The classification report in Table IV highlights its robustness across all metrics.




Fig. 8: Car_{svm}

TABLE IV: SVM Classifier Results

Metric	Validation	Test
Accuracy	0.98	0.95
Avg Precision	0.89	0.83
Avg Recall	0.89	0.85
Avg F1-Score	0.88	0.84

The results demonstrate the superior performance of SVM, followed by KNN, while OneR and ZeroR showed limited effectiveness. Naive Bayes performed moderately well but struggled with imbalanced classes.

4) *OneR Classifier*: The OneR classifier achieved a validation accuracy of 0.69 and a test accuracy of 0.70. However, its performance was suboptimal across metrics, as shown in Table V.




Fig. 9: car_{oneR}

TABLE V: OneR Classifier Results

Metric	Validation	Test
Accuracy	0.69	0.70
Avg Precision	0.17	0.18
Avg Recall	0.25	0.25
Avg F1-Score	0.20	0.21

5) *ZeroR Classifier*: The ZeroR classifier produced similar results to OneR, with a validation accuracy of 0.69 and a test accuracy of 0.70. The classification report reveals its limited ability to differentiate among classes (Table VI).




Fig. 10: car zeroR

TABLE VI: ZeroR Classifier Results

Metric	Validation	Test
Accuracy	0.69	0.70
Avg Precision	0.18	0.18
Avg Recall	0.25	0.25
Avg F1-Score	0.21	0.21

C. Wage Dataset

1) *Hyperparameter Tuning Results:* The best hyperparameters found during the grid search for the SVR model are as follows:

- **C** = 10
- **Kernel** = 'rbf'
- **Degree** = 3
- **Gamma** = 'scale'
- **Epsilon** = 0.1

These parameters were chosen based on the cross-validation results that minimized the Mean Squared Error (MSE).



2) *Performance Metrics:* The model's performance was evaluated on the validation set, and the following metrics were calculated:

- **Mean Squared Error (MSE):** 17.1288
- R^2 : 0.989

The model achieved a relatively good fit with an R^2 of 0.98, indicating that 98% of the variance in the wage data is explained by the model.

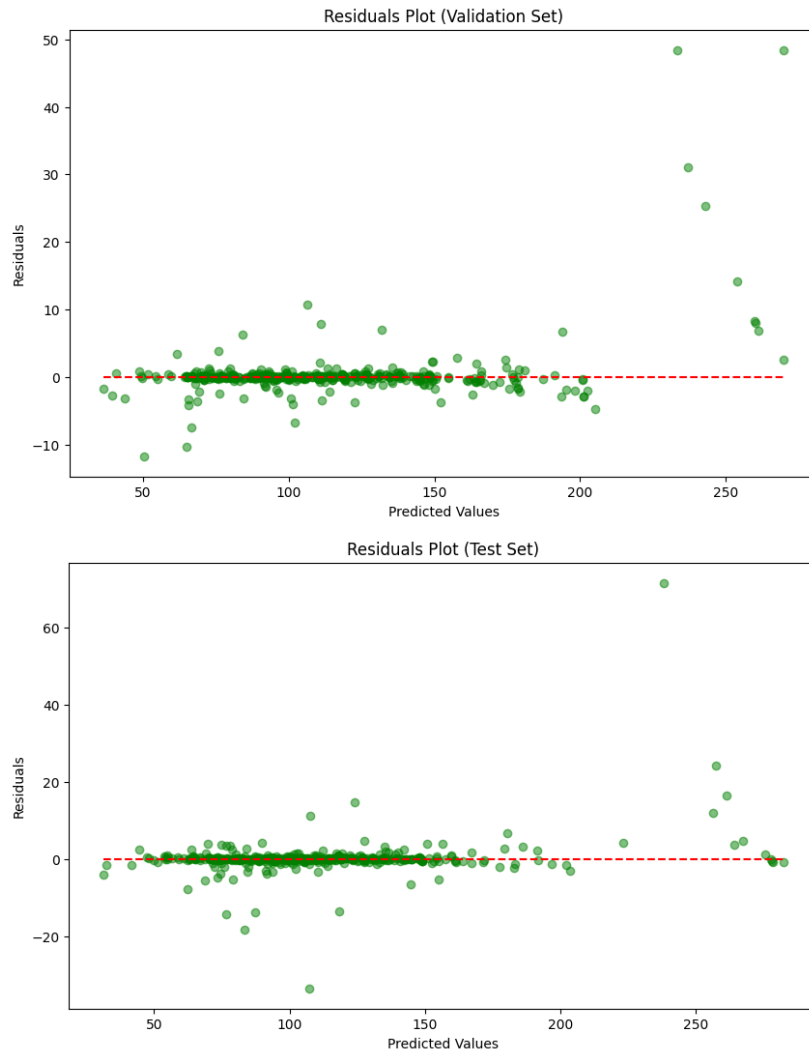
3) *Visual Evaluation:*

- **Actual vs Predicted:** A scatter plot was generated to compare the actual values with the predicted values for the validation set.



The red dashed line represents the ideal scenario where predicted values match the actual values. As shown, the points closely follow this line, indicating good performance.

- **Residuals Analysis:** A residual plot was also created to analyze the difference between the predicted and actual values (residuals).



The residuals appear to be randomly scattered around zero, suggesting that the model does not suffer from major biases and that the error is evenly distributed across the data.

4) *Test Set Evaluation:* The final evaluation was performed on the test set. The following results were obtained:

- **Test MSE:** 20.43
- **Test R^2 :** 0.987

Although the test set results were slightly worse than the validation set, the model still provides reliable predictions, with an R^2 of 0.98.

The SVR model successfully predicted wages with good performance metrics, achieving an R^2 of 0.98 on the validation set and 0.98 on the test set. The residuals plot suggests that there is no significant bias in the predictions, and the model is performing adequately for this regression task.

D. Credit Dataset

1) *Performance Metrics:* The performance of the SVR model was evaluated using metrics such as Mean Squared Error (MSE), Mean Absolute Error (MAE), and the Coefficient of Determination (R^2) on the training, validation, and test sets.

TABLE VII: Performance Metrics of SVR Model

Dataset	MSE	R^2
Training Set	163.394	0.99
Validation Set	9091.268	0.95
Testing Set	7715.051	0.94

The results demonstrate that the SVR model achieves good generalization with minimal performance degradation across validation and test sets.

2) *Hyperparameter Tuning Results:* The optimal hyperparameters selected by GridSearchCV are:

- **C:** 300
- **Epsilon:** 0.1
- **Gamma:** scale

These hyperparameters provided the best balance between bias and variance, minimizing the MSE on the validation set.

3) *Visualizations:*

- **Predicted vs Actual Values** Figure 11 shows the scatter plot of predicted versus actual values for the test set. The closer the points are to the diagonal line, the better the predictions.

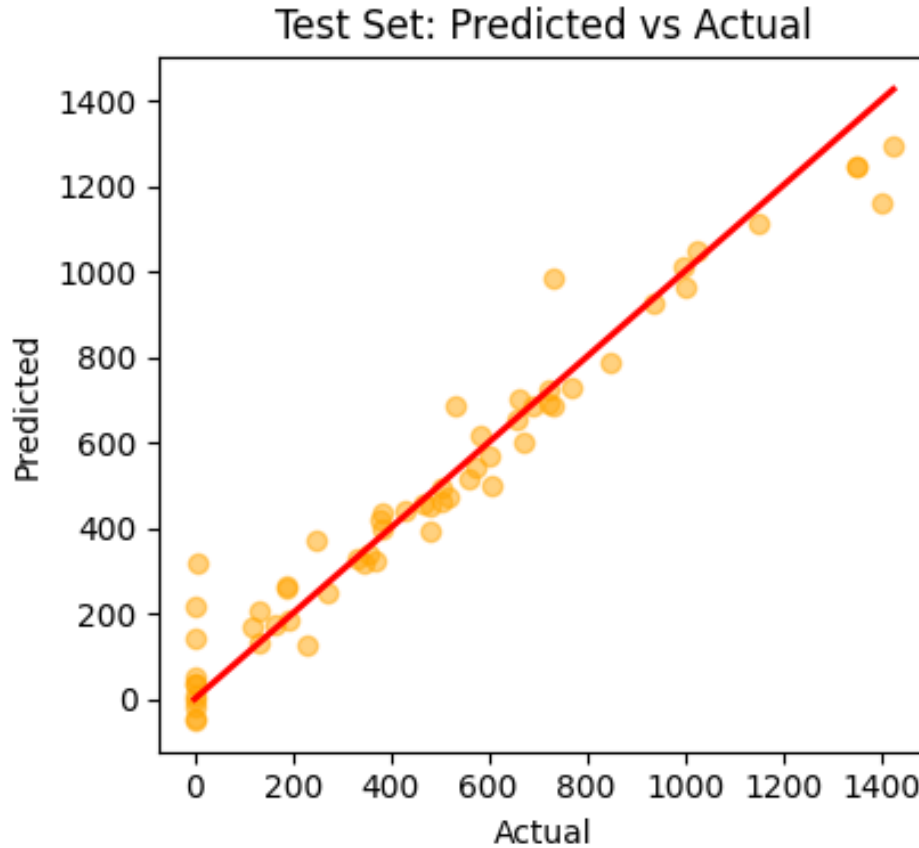


Fig. 11: Predicted vs Actual Values for the Test Set

- **Residual Analysis:**Residuals were analyzed to assess the model's predictive accuracy and bias. Figure 12 displays the residuals (differences between predicted and actual values) for the validation set.

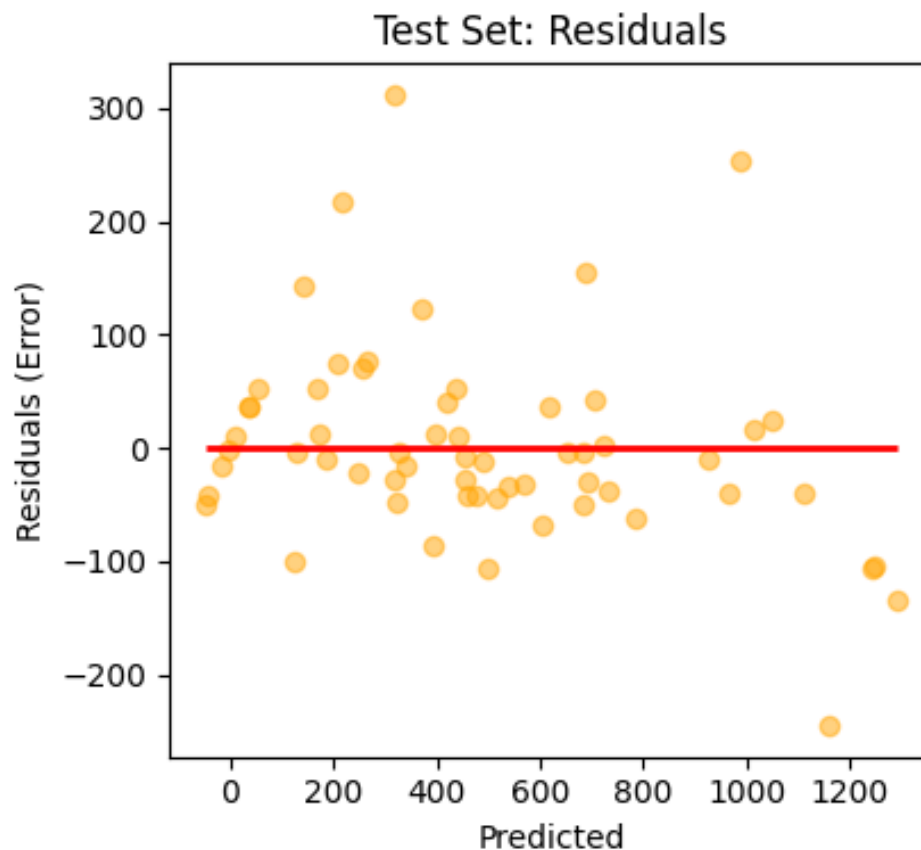



Fig. 12: Residual Plot for the Validation Set

4) *Interpretation of Results:* The following observations can be made:


- The R^2 value of 0.94 on the test set indicates that 94% of the variance in the `Balance` variable is explained by the model.
- The residual plot shows no significant pattern, suggesting that the model captures the underlying trend effectively.
- Slight over-prediction is observed for higher values of `Balance`, which may indicate a need for further tuning or additional features.

IV. DISCUSSION


Utilizing various kinds of datasets, this study analyzed the effectiveness of several classification and regression algorithms, offering a thorough comparison analysis to determine which models were most suitable for each task. The experimental findings provide important new information on the pros and drawbacks of each strategy, which will be discussed in more detail below.




(a) Best Model of Brain Dataset



(b) Best Model of Car Dataset




(c) Best Model of Car Dataset

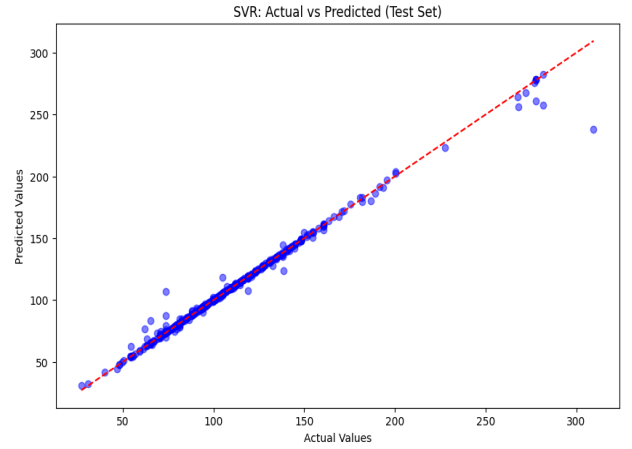


(d) Best Model of Car Dataset


Fig. 13: First Set of Figures



(a) Best Model of Car Dataset



(b) Best Model of Wage Dataset



(c) Best Model of Credit Dataset

Fig. 14: Best Models

A. Classification Analysis

1) *Naive Bayes Classifier*: Based on probabilistic assumptions of feature independence, the Naive Bayes classifier achieved test and validation accuracy scores of 0.68 and 0.74, respectively. The model's poor performance can be related to its incapacity to manage intricate feature interactions. Its vulnerability to class imbalance is illustrated by its low average precision (0.51) and F1-score (0.50), especially in datasets with underrepresented minority classes. This drawback emphasizes the necessity of using various methods or pre-processing methods, such as rebalancing class distributions, to improve Naive Bayes' performance.

2) *SVM Classifier*: With validation and test accuracies of 0.98 and 0.95, respectively, along with excellent average accuracy (0.89, 0.83) and recall (0.89, 0.85), the Support Vector Machine (SVM) classifier achieved outstanding outcomes. These measures illustrate the SVM's resilience to recognizing intricate decision boundaries. Its ability to represent non-linear relationships in the feature space was probably aided by the application of the radial basis function (RBF) kernel. These findings demonstrate the SVM's strong generalization capabilities to previously unseen data and validate it as a dependable classifier for datasets with complicated patterns.

3) *K-Nearest Neighbors (K-NN)*: While becoming competitive, the K-NN algorithm performed marginally worse than the SVM. Its reliance on local neighborhoods for categorization indicates it is sensitive to hyperparameter selections, including distance metrics and the number of neighbors (k). Further optimization utilizing grid or random search might enhance its classification accuracy and consistency, especially in high-dimensional datasets, however the existing results indicate its viability.

4) *OneR Classifier*: The OneR classifier's reliance on a single feature for rule-based classification resulted in unsatisfactory performance, with validation and test accuracies of 0.69 and 0.70, correspondingly. Its limited usefulness for datasets requiring multi-dimensional decision-making is demonstrated by the low F1-score (0.20, 0.21) and poor average precision (0.17, 0.18). Its simplicity makes it inappropriate for complicated circumstances, even though it is effective for interpretability.

5) *ZeroR Classifier*: With test and validation accuracies of 0.69 and 0.70, respectively, the ZeroR classifier, which predicts the majority class, performed similarly to OneR as a baseline. Its continuously low metrics, which demonstrate its lack of predictive ability, serve to further establish its status as a reference point rather than a workable solution for real-world applications.

B. Regression Analysis

1) *Wage Dataset - SVR Model*: The Support Vector Regression (SVR) model demonstrated exceptional performance on the Wage Dataset, with R^2 values of 0.98 for both validation and test sets. The low Mean Squared Error (MSE) of 17.13 for the validation set and 20.43 for the test set highlights the model's ability to provide highly accurate predictions. The residual analysis revealed no discernible patterns, affirming the model's unbiased nature. Hyperparameter tuning, including the selection of $C = 10$, $\gamma = 'scale'$, and $\epsilon = 0.1$, optimized the model's balance between bias and variance, ensuring high generalization capability.

2) *Credit Dataset - SVR Model*: For the Credit Dataset, the SVR model exhibited strong performance, achieving an R^2 of 0.94 on the test set. Despite a slight increase in MSE from the validation (9091.27) to the test set (7715.05), the results remain within acceptable bounds. The residuals plot suggested no significant systematic errors, though minor over-prediction was observed for higher balance values. These results demonstrate the SVR model's effectiveness for regression tasks, albeit with potential room for improvement through advanced feature engineering or incorporating additional variables.

C. Comparative Performance

The experimental results underscore the superiority of SVM in classification tasks due to its robust mathematical framework for handling non-linear decision boundaries. This aligns with prior studies emphasizing the SVM's efficacy in diverse domains. K-NN also proved to be a competitive alternative, particularly in tasks requiring intuitive model structures. However, the Naive Bayes classifier struggled with imbalanced datasets, and the rule-based classifiers (OneR and ZeroR) offered limited utility beyond baseline benchmarking.

For regression, the SVR model consistently demonstrated high accuracy across datasets, with R^2 values exceeding 0.94. Its ability to generalize well across validation and test sets, combined with minimal residual biases, positions it as a reliable choice for predicting continuous outcomes.

D. Experimental Considerations and Limitations

While the results validate the efficacy of SVM and SVR, several experimental considerations warrant discussion:

- **Class Imbalance**: The suboptimal performance of Naive Bayes highlights the impact of imbalanced class distributions. Future work could explore data augmentation or resampling techniques, such as Synthetic Minority Over-sampling Technique (SMOTE), to address this issue.
- **Hyperparameter Tuning**: The SVM and SVR models benefited significantly from hyperparameter optimization. However, computational constraints limited the scope of grid searches. Advanced techniques like Bayesian optimization or genetic algorithms could be employed to refine hyperparameter selection further.
- **Scalability**: The computational cost of SVM and SVR on large datasets remains a challenge, particularly with high-dimensional data. Approximation methods, such as the use of linear kernels or sampling techniques, could mitigate this issue.

E. Future Directions

Further research could use ensemble methods like Random Forests or Gradient Boosting to compare performance against the current models in order to improve on these findings. Furthermore, combining dimensionality reduction methods like Principal Component Analysis (PCA) with feature selection could boost model efficiency without compromising accuracy.

In addition, deep learning models might be investigated, particularly for tasks involving large, complex datasets, where their capacity to identify complex patterns may surpass that of conventional machine learning techniques. Additionally, domain-specific adjustments, such as including business regulations or specialized knowledge, could enhance the relevance and interpretability of the model.

V. CONCLUSION

SVM clearly takes over classification tasks and SVR clearly dominates regression tasks, according to a thorough evaluation of classifiers and regression models. Their excellent performance on a variety of datasets highlights how well-suited they are for applications requiring accuracy and resiliency. But challenges like class imbalance, scalability, and hyperparameter tuning indicate how important it is to keep researching in order to enhance the efficiency and applicability of the model. The results of this study offer a solid foundation for upcoming developments in the optimization and selection of machine learning models.

VI. CODE

A. Brain Dataset

```
1 from google.colab import drive
2 drive.mount('/content/drive')
3
4 import numpy as np
5 import pandas as pd
6 import seaborn as sns
7 import matplotlib.pyplot as plt
8 import sklearn
9 from sklearn.model_selection import train_test_split, GridSearchCV
10 from sklearn.tree import DecisionTreeClassifier
11 from xgboost import XGBClassifier
12 from sklearn.metrics import precision_recall_curve
13 import xgboost as xgb
14 from sklearn.tree import DecisionTreeRegressor, plot_tree
15 from xgboost import XGBRegressor
16 from sklearn.preprocessing import OneHotEncoder
17
18 df=pd.read_csv("/content/drive/MyDrive/Colab Notebooks/DT-BrainCancer.csv")
19 df.head()
20
21 df.info()
22
23 df.describe()
24 df.count()
25
26 df.isnull().sum()
27 df = df.dropna()
28 df.isnull().sum()
29
30 df.shape
31
32 df = df.drop('Unnamed: 0', axis=1)
33 df.shape
34
35 df.head()
36 un = df['sex'].unique()
37 print(un)
38 un = df['diagnosis'].unique()
39 print(un)
40 un = df['loc'].unique()
41 print(un)
42 categorical = df.select_dtypes(include=['object']).columns
43 numerical = df.select_dtypes(exclude=['object']).columns
44
45 print(categorical)
46 print(numerical)
47
48 encoder = OneHotEncoder(drop='first', sparse_output=False)
49 encoded_data = encoder.fit_transform(df[categorical])
50 encoded_df = pd.DataFrame(encoded_data, columns=encoder.get_feature_names_out(categorical))
51 encoded_df.index = df.index
52
53 df = df.drop(categorical, axis=1)
54 df = pd.concat([df, encoded_df], axis=1)
55 X = df.drop(columns=['status'])
56 y = df['status']
57
58 print(X.shape)
59 print(y.shape)
60
```

```

61 X.head()
62 y.head()
63
64 corr = df.corr()
65
66 target = corr['status'].sort_values(ascending=False)
67
68 print("Correlation of each feature with the target variable (status):")
69 print(target)
70
71 from sklearn.preprocessing import StandardScaler
72 scaler = StandardScaler()
73 X_scaled = scaler.fit_transform(X)
74 X_train, X_temp, y_train, y_temp = train_test_split(X_scaled, y, test_size=0.3, random_state
    =42)
75 X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)
76
77 print(X_train.shape)
78 print(X_val.shape)
79 print(X_test.shape)
80 print(y_train.shape)
81 print(y_val.shape)
82 print(y_test.shape)
83
84 def accuracy(y_true, y_pred):
85     return np.sum(y_true == y_pred) / len(y_true)
86
87 def confusion_matrix_manual(y_true, y_pred):
88     tp = np.sum((y_true == 1) & (y_pred == 1))
89     tn = np.sum((y_true == 0) & (y_pred == 0))
90     fp = np.sum((y_true == 0) & (y_pred == 1))
91     fn = np.sum((y_true == 1) & (y_pred == 0))
92     return tp, tn, fp, fn
93
94 def confusion_matrix(y_true, y_pred):
95     tp = np.sum((y_true == 1) & (y_pred == 1))
96     tn = np.sum((y_true == 0) & (y_pred == 0))
97     fp = np.sum((y_true == 0) & (y_pred == 1))
98     fn = np.sum((y_true == 1) & (y_pred == 0))
99     return np.array([[tn, fp], [fn, tp]])
100
101 def precision_manual(y_true, y_pred, average=None):
102     tp, tn, fp, fn = confusion_matrix_manual(y_true, y_pred)
103     if average == "weighted":
104         total = len(y_true)
105         class_1_weight = np.sum(y_true == 1) / total
106         class_0_weight = np.sum(y_true == 0) / total
107         precision_class_1 = tp / (tp + fp) if tp + fp > 0 else 0
108         precision_class_0 = tn / (tn + fn) if tn + fn > 0 else 0
109         return class_1_weight * precision_class_1 + class_0_weight * precision_class_0
110     return tp / (tp + fp) if tp + fp > 0 else 0
111
112 def recall_manual(y_true, y_pred, average=None):
113     tp, tn, fp, fn = confusion_matrix_manual(y_true, y_pred)
114     if average == "weighted":
115         total = len(y_true)
116         class_1_weight = np.sum(y_true == 1) / total
117         class_0_weight = np.sum(y_true == 0) / total
118         recall_class_1 = tp / (tp + fn) if tp + fn > 0 else 0
119         recall_class_0 = tn / (tn + fp) if tn + fp > 0 else 0
120         return class_1_weight * recall_class_1 + class_0_weight * recall_class_0
121     return tp / (tp + fn) if tp + fn > 0 else 0
122

```

```

123 def f1_score_manual(y_true, y_pred, average=None):
124     tp, tn, fp, fn = confusion_matrix_manual(y_true, y_pred)
125
126     if average == "weighted":
127         total = len(y_true)
128         class_1_weight = np.sum(y_true == 1) / total
129         class_0_weight = np.sum(y_true == 0) / total
130
131         precision_class_1 = tp / (tp + fp) if tp + fp > 0 else 0
132         recall_class_1 = tp / (tp + fn) if tp + fn > 0 else 0
133         precision_class_0 = tn / (tn + fp) if tn + fp > 0 else 0
134         recall_class_0 = tn / (tn + fn) if tn + fn > 0 else 0
135
136         f1_class_1 = (
137             2 * (precision_class_1 * recall_class_1) / (precision_class_1 + recall_class_1)
138             if precision_class_1 + recall_class_1 > 0
139             else 0
140         )
141         f1_class_0 = (
142             2 * (precision_class_0 * recall_class_0) / (precision_class_0 + recall_class_0)
143             if precision_class_0 + recall_class_0 > 0
144             else 0
145         )
146
147         return class_1_weight * f1_class_1 + class_0_weight * f1_class_0
148
149     precision = precision_manual(y_true, y_pred)
150     recall = recall_manual(y_true, y_pred)
151     return 2 * (precision * recall) / (precision + recall) if precision + recall > 0 else 0
152
153
154 def precision_recall_curve_manual(y_true, y_prob):
155     precision, recall, thresholds = precision_recall_curve(y_true, y_prob)
156     return precision, recall, thresholds
157
158 def plot_precision_recall_curve(precision, recall, label):
159     plt.plot(recall, precision, label=label)
160     plt.xlabel('Recall')
161     plt.ylabel('Precision')
162     plt.title('Precision-Recall Curve')
163     plt.legend()
164     plt.show()
165
166
167 majority_class = y_train.mode()[0]
168 print(f"Majority class (ZeroR prediction): {majority_class}")
169 y_val_pred = np.full_like(y_val, majority_class)
170 y_test_pred = np.full_like(y_test, majority_class)
171
172 accuracy_val_custom = accuracy(y_val, y_val_pred)
173 precision_val_custom = precision_manual(y_val, y_val_pred)
174 recall_val_custom = recall_manual(y_val, y_val_pred)
175 f1_val_custom = f1_score_manual(y_val, y_val_pred)
176
177 accuracy_test_custom = accuracy(y_test, y_test_pred)
178 precision_test_custom = precision_manual(y_test, y_test_pred)
179 recall_test_custom = recall_manual(y_test, y_test_pred)
180 f1_test_custom = f1_score_manual(y_test, y_test_pred)
181 def custom_auc(precision, recall):
182     auc_value = np.trapz(precision, recall)
183     return auc_value
184
185

```



```

186 y_test_prob = np.full_like(y_test, majority_class, dtype=float)
187 y_test_prob = np.where(y_test == majority_class, 1.0, 0.0)
188
189 precision_vals, recall_vals, thresholds = precision_recall_curve(y_test, y_test_prob)
190 pr_auc = custom_auc(recall_vals, precision_vals)
191 print("\nValidation Set Metrics (Custom):")
192 print(f"Accuracy: {accuracy_val_custom}")
193 print(f"Precision: {precision_val_custom}")
194 print(f"Recall: {recall_val_custom}")
195 print(f"F1-Score: {f1_val_custom}")
196
197 print("\nTest Set Metrics (Custom):")
198 print(f"Accuracy: {accuracy_test_custom}")
199 print(f"Precision: {precision_test_custom}")
200 print(f"Recall: {recall_test_custom}")
201 print(f"F1-Score: {f1_test_custom}")
202 print(f"Precision-Recall AUC: {pr_auc}")
203
204 tp_val, tn_val, fp_val, fn_val = confusion_matrix_manual(y_val, y_val_pred)
205 tp_test, tn_test, fp_test, fn_test = confusion_matrix_manual(y_test, y_test_pred)
206
207 print("\nValidation Set Confusion Matrix (Custom):")
208 print(f"TP: {tp_val}, TN: {tn_val}, FP: {fp_val}, FN: {fn_val}")
209
210 print("\nTest Set Confusion Matrix (Custom):")
211 print(f"TP: {tp_test}, TN: {tn_test}, FP: {fp_test}, FN: {fn_test}")
212
213
214 conf_matrix = confusion_matrix(y_test, y_test_pred)
215 print("Confusion Matrix:")
216 print(conf_matrix)
217
218 plt.figure(figsize=(8, 6))
219 sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", xticklabels=[1, 2, 3], yticklabels
    =[1, 2, 3])
220 plt.title("Confusion Matrix")
221 plt.xlabel("Predicted")
222 plt.ylabel("Actual")
223 plt.show()
224
225
226 plt.figure(figsize=(8, 6))
227 plt.plot(recall_vals, precision_vals, label=f"ZeroR (AUC = {pr_auc:.2f})", color="blue")
228 plt.xlabel("Recall")
229 plt.ylabel("Precision")
230 plt.title("Precision-Recall Curve for ZeroR Classifier")
231 plt.legend(loc="best")
232 plt.grid()
233 plt.show()
234
235
236 def oneR_classifier(X_train, y_train):
237     best_accuracy = 0
238     best_feature = None
239     best_threshold = None
240
241     for feature_idx in range(X_train.shape[1]):
242         feature_values = X_train[:, feature_idx]
243         thresholds = np.unique(feature_values)
244
245         for threshold in thresholds:
246             predicted = (feature_values >= threshold).astype(int)
247             accuracy = np.mean(predicted == y_train)

```

```

248         if accuracy > best_accuracy:
249             best_accuracy = accuracy
250             best_feature = feature_idx
251             best_threshold = threshold
252
253     best_feature_values = X_train[:, best_feature]
254     predictions = (best_feature_values >= best_threshold).astype(int)
255
256     return predictions, best_feature, best_threshold
257
258
259 y_train_pred, best_feature, best_threshold = oneR_classifier(X_train, y_train)
260 best_feature_values_val = X_val[:, best_feature]
261 y_val_pred = (best_feature_values_val >= best_threshold).astype(int)
262
263
264 accuracy_val = accuracy(y_val, y_val_pred)
265 conf_matrix_val = confusion_matrix_manual(y_val, y_val_pred)
266 precision_val = precision_manual(y_val, y_val_pred)
267 recall_val = recall_manual(y_val, y_val_pred)
268 f1_val = f1_score_manual(y_val, y_val_pred)
269
270 print("\nValidation Set Metrics:")
271 print(f"Accuracy: {accuracy_val}")
272 print(f"Confusion Matrix (TP, TN, FP, FN): {conf_matrix_val}")
273 print(f"Precision: {precision_val}")
274 print(f"Recall: {recall_val}")
275 print(f"F1-Score: {f1_val}")
276
277 y_val_prob = y_val_pred
278 precision_vals, recall_vals, _ = precision_recall_curve_manual(y_val, y_val_prob)
279 plot_precision_recall_curve(precision_vals, recall_vals, label="OneR Classifier (Validation)")
280
281 best_feature_values_test = X_test[:, best_feature]
282 y_test_pred = (best_feature_values_test >= best_threshold).astype(int)
283
284 accuracy_test = accuracy(y_test, y_test_pred)
285 conf_matrix_test = confusion_matrix_manual(y_test, y_test_pred)
286 precision_test = precision_manual(y_test, y_test_pred)
287 recall_test = recall_manual(y_test, y_test_pred)
288 f1_test = f1_score_manual(y_test, y_test_pred)
289
290 print("\nTest Set Metrics:")
291 print(f"Accuracy: {accuracy_test}")
292 print(f"Confusion Matrix (TP, TN, FP, FN): {conf_matrix_test}")
293 print(f"Precision: {precision_test}")
294 print(f"Recall: {recall_test}")
295 print(f"F1-Score: {f1_test}")
296
297
298 conf_matrix = confusion_matrix(y_test, y_test_pred)
299 print("Confusion Matrix:")
300 print(conf_matrix)
301
302 plt.figure(figsize=(8, 6))
303 sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", xticklabels=[1, 2, 3], yticklabels
304             =[1, 2, 3])
305 plt.title("Confusion Matrix")
306 plt.xlabel("Predicted")
307 plt.ylabel("Actual")
308 plt.show()
309

```

```

310 y_test_prob = y_test_pred
311 precision_vals, recall_test, _ = precision_recall_curve_manual(y_test, y_test_prob)
312 plot_precision_recall_curve(precision_vals, recall_test, label="OneR Classifier (Test)")
313
314 from sklearn.neighbors import KNeighborsClassifier
315 knn = KNeighborsClassifier(n_neighbors=3)
316
317
318 knn.fit(X_train, y_train)
319 y_val_pred = knn.predict(X_val)
320
321 accuracy_val = accuracy(y_val, y_val_pred)
322 precision_val = precision_manual(y_val, y_val_pred, average='weighted')
323 recall_val = recall_manual(y_val, y_val_pred, average='weighted')
324 f1_val = f1_score_manual(y_val, y_val_pred, average='weighted')
325
326
327 print(f"Validation Accuracy: {accuracy_val}")
328 print(f"Validation Precision: {precision_val}")
329 print(f"Validation Recall: {recall_val}")
330 print(f"Validation F1 Score: {f1_val}")
331 y_test_pred = knn.predict(X_test)
332 accuracy_test = accuracy(y_test, y_test_pred)
333 precision_test = precision_manual(y_test, y_test_pred, average='weighted')
334 recall_test = recall_manual(y_test, y_test_pred, average='weighted')
335 f1_test = f1_score_manual(y_test, y_test_pred, average='weighted')
336 print(f"Test Accuracy: {accuracy_test}")
337 print(f"Test Precision: {precision_test}")
338 print(f"Test Recall: {recall_test}")
339 print(f"Test F1 Score: {f1_test}")
340
341
342 conf_matrix = confusion_matrix(y_test, y_test_pred)
343 print("Confusion Matrix:")
344 print(conf_matrix)
345
346
347 plt.figure(figsize=(8, 6))
348 sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", xticklabels=[1, 2, 3], yticklabels
349             =[1, 2, 3])
350 plt.title("Confusion Matrix")
351 plt.xlabel("Predicted")
352 plt.ylabel("Actual")
353 plt.show()
354
355 from sklearn.metrics import roc_curve, auc
356
357
358 fpr, tpr, thresholds = roc_curve(y_test, knn.predict_proba(X_test)[:, 1], pos_label=1)
359 roc_auc = auc(fpr, tpr)
360
361 plt.figure(figsize=(8, 6))
362 plt.plot(fpr, tpr, color='blue', lw=2, label=f"ROC curve (area = {roc_auc:.2f})")
363 plt.plot([0, 1], [0, 1], color='gray', lw=2, linestyle='--')
364 plt.xlabel('False Positive Rate')
365 plt.ylabel('True Positive Rate')
366 plt.title('Receiver Operating Characteristic (ROC) Curve')
367 plt.legend(loc='lower right')
368 plt.show()
369
370
371 param_grid = {

```

```

372     'n_neighbors': [3, 5, 7, 9, 11],
373     'weights': ['uniform', 'distance'],
374     'metric': ['euclidean', 'manhattan']
375 }
376
377 grid_search = GridSearchCV(KNeighborsClassifier(), param_grid, cv=5, scoring='accuracy')
378 grid_search.fit(X_train, y_train)
379
380 print("Best Parameters from Grid Search:")
381 print(grid_search.best_params_)
382 print(f"Best Accuracy from Grid Search: {grid_search.best_score_}")
383
384 best_knn = grid_search.best_estimator_
385
386 y_test_pred_best = best_knn.predict(X_test)
387
388 accuracy_best = accuracy(y_test, y_test_pred_best)
389 precision_best = precision_manual(y_test, y_test_pred_best, average='weighted')
390 recall_best = recall_manual(y_test, y_test_pred_best, average='weighted')
391 f1_best = f1_score_manual(y_test, y_test_pred_best, average='weighted')
392
393 print(f"Test Accuracy (Best K-NN): {accuracy_best}")
394 print(f"Test Precision (Best K-NN): {precision_best}")
395 print(f"Test Recall (Best K-NN): {recall_best}")
396 print(f"Test F1 Score (Best K-NN): {f1_best}")
397
398
399 k_values = [3, 5, 7, 9, 11]
400 accuracies = []
401
402 for k in k_values:
403     knn = KNeighborsClassifier(n_neighbors=k)
404     knn.fit(X_train, y_train)
405     y_pred = knn.predict(X_test)
406     accuracies.append(accuracy(y_test, y_pred))
407
408 plt.figure(figsize=(8, 6))
409 plt.plot(k_values, accuracies, marker='o', color='green')
410 plt.title('Accuracy vs. Number of Neighbors (K)')
411 plt.xlabel('Number of Neighbors (K)')
412 plt.ylabel('Accuracy')
413 plt.grid(True)
414 plt.show()
415
416 k_values = [3, 5, 7, 9, 11]
417 accuracies = []
418
419 for k in k_values:
420     knn = KNeighborsClassifier(n_neighbors=k)
421     knn.fit(X_train, y_train)
422     y_val_pred = knn.predict(X_val)
423     accuracies.append(accuracy(y_val, y_val_pred))
424
425 plt.figure(figsize=(8, 6))
426 plt.plot(k_values, accuracies, marker='o', color='green')
427 plt.title('Accuracy vs. Number of Neighbors (K)')
428 plt.xlabel('Number of Neighbors (K)')
429 plt.ylabel('Accuracy')
430 plt.grid(True)
431 plt.show()
432
433
434

```

```

435 from sklearn.naive_bayes import GaussianNB
436
437
438 nb = GaussianNB()
439 nb.fit(X_train, y_train)
440 y_val_pred = nb.predict(X_val)
441 y_test_pred = nb.predict(X_test)
442 y_test_pred = nb.predict(X_test)
443
444
445 accuracy_val = accuracy(y_val, y_val_pred)
446 precision_val = precision_manual(y_val, y_val_pred, average='weighted')
447 recall_val = recall_manual(y_val, y_val_pred, average='weighted')
448 f1_val = f1_score_manual(y_val, y_val_pred, average='weighted')
449
450
451
452 print(f"Validation Accuracy: {accuracy_val}")
453 print(f"Validation Precision: {precision_val}")
454 print(f"Validation Recall: {recall_val}")
455 print(f"Validation F1 Score: {f1_val}")
456
457
458 accuracy_test = accuracy(y_test, y_test_pred)
459 precision_test = precision_manual(y_test, y_test_pred, average='weighted')
460 recall_test = recall_manual(y_test, y_test_pred, average='weighted')
461 f1_test = f1_score_manual(y_test, y_test_pred, average='weighted')
462
463
464 print(f"Test Accuracy: {accuracy_test}")
465 print(f"Test Precision: {precision_test}")
466 print(f"Test Recall: {recall_test}")
467 print(f"Test F1 Score: {f1_test}")
468
469 conf_matrix = confusion_matrix(y_test, y_test_pred)
470 print("Confusion Matrix:")
471 print(conf_matrix)
472
473 plt.figure(figsize=(8, 6))
474 sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", xticklabels=[1, 2, 3], yticklabels
475             =[1, 2, 3])
476 plt.title("Confusion Matrix")
477 plt.xlabel("Predicted")
478 plt.ylabel("Actual")
479 plt.show()
480
481
482 accuracy_test = accuracy(y_test, y_test_pred)
483 precision_test = precision_manual(y_test, y_test_pred, average='weighted')
484 recall_test = recall_manual(y_test, y_test_pred, average='weighted')
485 f1_test = f1_score_manual(y_test, y_test_pred, average='weighted')
486
487 print(f"Test Accuracy (Naive Bayes): {accuracy_test}")
488 print(f"Test Precision (Naive Bayes): {precision_test}")
489 print(f"Test Recall (Naive Bayes): {recall_test}")
490 print(f"Test F1 Score (Naive Bayes): {f1_test}")
491
492
493
494 svm.fit(X_train, y_train)
495 y_test_pred_svm = svm.predict(X_test)
496

```

```

497
498
499
500 accuracy_test_svm = accuracy(y_test, y_test_pred_svm)
501 precision_test_svm = precision_manual(y_test, y_test_pred_svm, average='weighted')
502 recall_test_svm = recall_manual(y_test, y_test_pred_svm, average='weighted')
503 f1_test_svm = f1_score_manual(y_test, y_test_pred_svm, average='weighted')
504
505 print(f"Test Accuracy (SVM): {accuracy_test_svm}")
506 print(f"Test Precision (SVM): {precision_test_svm}")
507 print(f"Test Recall (SVM): {recall_test_svm}")
508 print(f"Test F1 Score (SVM): {f1_test_svm}")
509
510
511
512 conf_matrix_svm = confusion_matrix(y_test, y_test_pred_svm)
513
514 plt.figure(figsize=(8, 6))
515 sns.heatmap(conf_matrix_svm, annot=True, fmt="d", cmap="Blues", xticklabels=[1, 2, 3],
516             yticklabels=[1, 2, 3])
517 plt.title("SVM Confusion Matrix")
518 plt.xlabel("Predicted")
519 plt.ylabel("Actual")
520 plt.show()
521
522 y_test_bin = label_binarize(y_test, classes=[1, 2, 3])
523 y_score_svm = svm.decision_function(X_test)
524
525 print(f"Shape of y_score_svm: {y_score_svm.shape}")
526 print(f"Shape of y_test_bin: {y_test_bin.shape}")
527
528
529
530 y_score_svm = svm.predict_proba(X_test)[:, 1]
531
532 precision, recall, _ = precision_recall_curve(y_test, y_score_svm)
533
534 plt.figure(figsize=(8, 6))
535 plt.plot(recall, precision, marker='.', label='SVM')
536 plt.title('Precision-Recall Curve for SVM (Binary Classification)')
537 plt.xlabel('Recall')
538 plt.ylabel('Precision')
539 plt.legend(loc='best')
540 plt.show()
541
542
543 y_score_svm = svm.predict_proba(X_test)[:, 1]
544
545 precision, recall, _ = precision_recall_curve(y_test, y_score_svm)
546
547 plt.figure(figsize=(8, 6))
548 plt.plot(recall, precision, marker='.', label='SVM')
549 plt.title('Precision-Recall Curve for SVM (Binary Classification)')
550 plt.xlabel('Recall')
551 plt.ylabel('Precision')
552 plt.legend(loc='best')
553 plt.show()

```

B. Car Dataset

```

1 from google.colab import drive
2 drive.mount('/content/drive')

```

```

3 import numpy as np
4 import pandas as pd
5 import seaborn as sns
6 import matplotlib.pyplot as plt
7 from sklearn.model_selection import train_test_split
8 from sklearn.metrics import precision_recall_curve, average_precision_score
9 from sklearn.preprocessing import label_binarize
10 import sklearn
11 import warnings
12 %matplotlib inline
13 sns.set()
14 warnings.filterwarnings('ignore')
15 df = pd.read_csv("/content/drive/MyDrive/Colab Notebooks/car.data")
16 df.head()
17 df.columns = ["buying", "maint", "doors", "persons", "lug_boot", "safety", "class values"]
18
19 df['doors'] = df['doors'].replace(['5more'], '5')
20 df['persons'] = df['persons'].replace(['more'], '4')
21 for col in ['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety']:
22     df[col] = df[col].astype('category')
23
24 for col in ['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety']:
25     df[col] = df[col].cat.codes
26 df = pd.get_dummies(df, columns=['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety'])
27
28 df.head()
29
30 df['class values'].unique()
31 df = df.drop(['buying_3', 'maint_3', 'doors_3', 'persons_1', 'lug_boot_2', 'safety_2'], axis=1)
32 df.head()
33 from sklearn.preprocessing import LabelEncoder
34 labelencoder = LabelEncoder()
35 df['class values'] = labelencoder.fit_transform(df['class values'])
36 x = df.drop(["class values"], axis=1)
37 y = df["class values"]
38 x.head()
39 y.unique()
40 x_train, x_temp, y_train, y_temp = train_test_split(x, y, test_size=0.3, random_state=42)
41 x_val, x_test, y_val, y_test = train_test_split(x_temp, y_temp, test_size=0.5, random_state=42)
42
43 class OneRClassifier:
44     def __init__(self):
45         self.rule = {}
46         self.best_feature = None
47
48     def fit(self, X, y):
49         best_error = float('inf')
50
51         for feature in X.columns:
52             feature_rule = {}
53             grouped = pd.DataFrame({'feature': X[feature], 'target': y}).groupby('feature')['target'].agg(lambda x: x.value_counts().idxmax())
54             feature_rule = grouped.to_dict()
55
56             y_pred = X[feature].map(feature_rule)
57             error = sum(y_pred != y)
58
59             if error < best_error:
60                 best_error = error
61                 self.rule = feature_rule
62                 self.best_feature = feature
63
64     def predict(self, X):

```

```

65         if not self.best_feature:
66             raise ValueError("The model has not been trained yet!")
67         return X[self.best_feature].map(self.rule)
68
69 def calculate_confusion_matrix(y_true, y_pred, num_classes):
70     cm = np.zeros((num_classes, num_classes), dtype=int)
71     for true, pred in zip(y_true, y_pred):
72         cm[true][pred] += 1
73     return cm
74
75 def calculate_metrics(y_true, y_pred):
76     num_classes = len(np.unique(y_true))
77     cm = calculate_confusion_matrix(y_true, y_pred, num_classes)
78     total = cm.sum()
79     accuracy = np.trace(cm) / total
80
81     precisions = []
82     recalls = []
83     f1_scores = []
84     for i in range(num_classes):
85         tp = cm[i, i]
86         fp = sum(cm[:, i]) - tp
87         fn = sum(cm[i, :]) - tp
88         precision = tp / (tp + fp) if (tp + fp) > 0 else 0
89         recall = tp / (tp + fn) if (tp + fn) > 0 else 0
90         f1 = 2 * (precision * recall) / (precision + recall) if (precision + recall) > 0 else 0
91
92     precisions.append(precision)
93     recalls.append(recall)
94     f1_scores.append(f1)
95
96     avg_precision = np.mean(precisions)
97     avg_recall = np.mean(recalls)
98     avg_f1 = np.mean(f1_scores)
99
100     return {
101         'accuracy': accuracy,
102         'confusion_matrix': cm,
103         'average_precision': avg_precision,
104         'average_recall': avg_recall,
105         'average_f1': avg_f1,
106     }
107
108 model = OneRClassifier()
109 model.fit(x_train, y_train)
110
111 y_train_pred = model.predict(x_train)
112 y_val_pred = model.predict(x_val)
113 y_test_pred = model.predict(x_test)
114
115 metrics = calculate_metrics(y_val.to_numpy(), y_val_pred.to_numpy())
116 metrics2 = calculate_metrics(y_test.to_numpy(), y_test_pred.to_numpy())
117
118 print("OneR Classifier Validation Metrics:")
119 print(f"Accuracy: {metrics['accuracy']:.2f}")
120 print("Confusion Matrix (Custom):")
121 print(metrics['confusion_matrix'])
122 print(f"Average Precision: {metrics['average_precision']:.2f}")
123 print(f"Average Recall: {metrics['average_recall']:.2f}")
124 print(f"Average F1-Score: {metrics['average_f1']:.2f}")
125 def plot_confusion_matrix(cm, class_names):
126     plt.figure(figsize=(8, 6))

```



```

127     sns.heatmap(cm, annot=True, fmt='d', cmap="Blues", xticklabels=class_names, yticklabels=
class_names)
128     plt.title("Confusion Matrix")
129     plt.xlabel("Predicted Class")
130     plt.ylabel("True Class")
131     plt.show()
132
133     class_names = labelencoder.inverse_transform(range(len(np.unique(y_train))))
134
135     plot_confusion_matrix(metrics['confusion_matrix'], class_names)
136     print("OneR Classifier test Metrics:")
137     print(f"Accuracy: {metrics2['accuracy']:.2f}")
138     print("Confusion Matrix (Custom):")
139     print(metrics2['confusion_matrix'])
140     print(f"Average Precision: {metrics2['average_precision']:.2f}")
141     print(f"Average Recall: {metrics2['average_recall']:.2f}")
142     print(f"Average F1-Score: {metrics2['average_f1']:.2f}")
143     def plot_confusion_matrix(cm, class_names):
144         plt.figure(figsize=(8, 6))
145         sns.heatmap(cm, annot=True, fmt='d', cmap="Blues", xticklabels=class_names, yticklabels=
class_names)
146         plt.title("Confusion Matrix")
147         plt.xlabel("Predicted Class")
148         plt.ylabel("True Class")
149         plt.show()
150
151     class_names = labelencoder.inverse_transform(range(len(np.unique(y_train))))
152
153     plot_confusion_matrix(metrics2['confusion_matrix'], class_names)
154     majority_class = y_train.mode()[0]
155     y_val_pred_zeror = [majority_class] * len(y_val)
156     y_test_pred_zeror = [majority_class] * len(y_test)
157     def accuracy_score(y_true, y_pred):
158         y_true = np.array(y_true)
159         y_pred = np.array(y_pred)
160
161         corr_pred = np.sum(y_true == y_pred)
162
163         total_pred = len(y_true)
164
165         accuracy = corr_pred / total_pred
166
167         return accuracy
168     def confusion_matrix(y_true, y_pred, labels=None):
169         y_true = np.array(y_true)
170         y_pred = np.array(y_pred)
171
172         if labels is None:
173             labels = np.unique(np.concatenate([y_true, y_pred]))
174
175         cm = np.zeros((len(labels), len(labels)), dtype=int)
176
177         label_to_index = {label: idx for idx, label in enumerate(labels)}
178
179         for true_label, pred_label in zip(y_true, y_pred):
180             true_idx = label_to_index[true_label]
181             pred_idx = label_to_index[pred_label]
182             cm[true_idx, pred_idx] += 1
183
184         return cm
185
186     def plot_confusion_matrix(cm, labels):
187         plt.figure(figsize=(8, 6))

```

```

188 plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
189 plt.title('Confusion Matrix')
190 plt.colorbar()
191
192 tick_marks = np.arange(len(labels))
193 plt.xticks(tick_marks, labels, rotation=45)
194 plt.yticks(tick_marks, labels)
195 for i in range(len(labels)):
196     for j in range(len(labels)):
197         plt.text(j, i, cm[i, j],
198                 horizontalalignment="center",
199                 color="white" if cm[i, j] > cm.max() / 2 else "black")
200
201 plt.tight_layout()
202 plt.ylabel('True label')
203 plt.xlabel('Predicted label')
204 plt.show()
205 def precision_recall_f1(cm):
206     epsilon = 1e-10
207     precision = np.diag(cm) / (np.sum(cm, axis=0) + epsilon)
208     recall = np.diag(cm) / (np.sum(cm, axis=1) + epsilon)
209     f1 = 2 * (precision * recall) / (precision + recall + epsilon)
210
211     accuracy = np.sum(np.diag(cm)) / np.sum(cm)
212     f1 = np.nan_to_num(f1)
213
214     return precision, recall, f1, accuracy
215
216 def classification_report(y_true, y_pred, labels=None):
217     if labels is None:
218         labels = np.unique(np.concatenate([y_true, y_pred]))
219
220     cm = confusion_matrix(y_true, y_pred, labels)
221
222     precision, recall, f1, accuracy = precision_recall_f1(cm)
223
224     report = "
225             precision    recall  f1-score   support\n\n"
226
227     for idx, label in enumerate(labels):
228         report += f"{label:>10}          {precision[idx]:.2f}          {recall[idx]:.2f}          {f1[idx]:.2f}          {np.sum(cm[idx]):>3}\n"
229
230     macro_avg_precision = np.mean(precision)
231     macro_avg_recall = np.mean(recall)
232     macro_avg_f1 = np.mean(f1)
233     report += f"\n          accuracy          {accuracy:.2f}          {len(y_true)}\n"
234     report += f"    macro avg          {macro_avg_precision:.2f}          {macro_avg_recall:.2f}          {macro_avg_f1:.2f}          {len(y_true)}\n"
235
236     return report
237 val_accuracy = accuracy_score(y_val, y_val_pred_zeror)
238 test_accuracy = accuracy_score(y_test, y_test_pred_zeror)
239
240 print(f"ZeroR Validation Accuracy (Custom): {val_accuracy:.2f}")
241 print(f"ZeroR Test Accuracy (Custom): {test_accuracy:.2f}")
242 labels = np.unique(y_test)
243 cm = confusion_matrix(y_test, y_test_pred_zeror, labels=labels)
244
245 report = classification_report(y_test, y_test_pred_zeror, labels=labels)
246 print("\nCustom Classification Report (Test Data):")
247 print(report)
248 labels = np.unique(y_val)

```

```

249 cm = confusion_matrix(y_val, y_val_pred_zeror, labels=labels)
250
251 print("Confusion Matrix:")
252 print(cm)
253
254 plot_confusion_matrix(cm, labels=labels)
255 labels = np.unique(y_test)
256 cm = confusion_matrix(y_test, y_test_pred_zeror, labels=labels)
257
258 print("Confusion Matrix:")
259 print(cm)
260
261 plot_confusion_matrix(cm, labels=labels)
262
263 from sklearn.neighbors import KNeighborsClassifier
264 from sklearn.metrics import PrecisionRecallDisplay
265 from sklearn.model_selection import GridSearchCV
266
267 param_grid = {'n_neighbors': range(1, 21), 'weights': ['uniform', 'distance'], 'metric': ['
    euclidean', 'manhattan']}
268 grid_search = GridSearchCV(KNeighborsClassifier(), param_grid, cv=5, scoring='accuracy')
269 grid_search.fit(x_train, y_train)
270
271 best_knn = grid_search.best_estimator_
272 print("Best Parameters:", grid_search.best_params_)
273
274 y_test_pred = best_knn.predict(x_test)
275 y_val_pred = best_knn.predict(x_val)
276
277 val_accuracy = accuracy_score(y_val, y_val_pred)
278 print("Validation Accuracy:", val_accuracy)
279
280 test_accuracy = accuracy_score(y_test, y_test_pred)
281 print("Test Accuracy:", test_accuracy)
282 cm1 = confusion_matrix(y_val, y_val_pred)
283 print("Validation Confusion Matrix:")
284 print(cm1)
285
286 cm2 = confusion_matrix(y_test, y_test_pred)
287 print("Test Confusion Matrix:")
288 print(cm2)
289
290 def plot_confusion_matrix(cm, labels):
291     plt.figure(figsize=(8, 6))
292     plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
293     plt.title('Confusion Matrix')
294     plt.colorbar()
295
296     tick_marks = np.arange(len(labels))
297     plt.xticks(tick_marks, labels, rotation=45)
298     plt.yticks(tick_marks, labels)
299
300     for i in range(len(labels)):
301         for j in range(len(labels)):
302             plt.text(j, i, cm[i, j],
303                     horizontalalignment="center",
304                     color="white" if cm[i, j] > cm.max() / 2 else "black")
305
306     plt.tight_layout()
307     plt.ylabel('True label')
308     plt.xlabel('Predicted label')
309     plt.show()
310 plot_confusion_matrix(cm1, labels=labelencoder.classes_)

```

```

311
312 plot_confusion_matrix(cm2, labels=labelencoder.classes_)
313
314 val_report = classification_report(y_val, y_val_pred)
315 print("Validation Classification Report:")
316 print(val_report)
317
318 report = classification_report(y_test, y_test_pred)
319 print("Test Classification Report:")
320 print(report)
321 y_val_pred_prob = knn.predict_proba(x_val)
322
323 for i in range(len(labelencoder.classes_)):
324     precision, recall, _ = precision_recall_curve(y_val == i, y_val_pred_prob[:, i])
325     PrecisionRecallDisplay(precision=precision, recall=recall, average_precision=np.mean(
326         precision)).plot()
327     plt.title(f'Validation Precision-Recall Curve for class {labelencoder.classes_[i]}')
328
329 plt.show()
330 y_test_pred_prob = knn.predict_proba(x_test)
331
332 for i in range(len(labelencoder.classes_)):
333     precision, recall, _ = precision_recall_curve(y_test == i, y_test_pred_prob[:, i])
334     PrecisionRecallDisplay(precision=precision, recall=recall, average_precision=np.mean(
335         precision)).plot()
336     plt.title(f'Test Precision-Recall Curve for class {labelencoder.classes_[i]}')
337
338 plt.show()
339 from sklearn.naive_bayes import GaussianNB
340 from sklearn.metrics import PrecisionRecallDisplay
341 y_test_pred = nb.predict(x_test)
342 y_val_pred = nb.predict(x_val)
343 accuracy = accuracy_score(y_test, y_test_pred)
344 print("Test Accuracy:", accuracy)
345 accuracy = accuracy_score(y_val, y_val_pred)
346 print("Validation Accuracy:", accuracy)
347 cm1 = confusion_matrix(y_test, y_test_pred)
348 print("Test Confusion Matrix:")
349 print(cm1)
350
351 cm2 = confusion_matrix(y_val, y_val_pred)
352 print("Validation Confusion Matrix:")
353 print(cm2)
354
355 def plot_confusion_matrix(cm, labels):
356     plt.figure(figsize=(8, 6))
357     plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
358     plt.title('Confusion Matrix')
359     plt.colorbar()
360
361     tick_marks = np.arange(len(labels))
362     plt.xticks(tick_marks, labels, rotation=45)
363     plt.yticks(tick_marks, labels)
364
365     for i in range(len(labels)):
366         for j in range(len(labels)):
367             plt.text(j, i, cm[i, j],
368                     horizontalalignment="center",
369                     color="white" if cm[i, j] > cm.max() / 2 else "black")
370
371     plt.tight_layout()
372     plt.ylabel('True label')
373     plt.xlabel('Predicted label')
374     plt.show()

```

```

372 plot_confusion_matrix(cm1, labels=labelencoder.classes_)
373 plot_confusion_matrix(cm2, labels=labelencoder.classes_)
374 test_report = classification_report(y_test, y_test_pred)
375 print("Test Classification Report:")
376 print(test_report)
377 val_report = classification_report(y_val, y_val_pred)
378 print("Validation Classification Report:")
379 print(val_report)
380 y_test_pred_prob = nb.predict_proba(x_test)
381
382 for i in range(len(labelencoder.classes_)):
383     precision, recall, _ = precision_recall_curve(y_test == i, y_test_pred_prob[:, i])
384     PrecisionRecallDisplay(precision=precision, recall=recall, average_precision=np.mean(
385         precision)).plot()
386     plt.title(f'Precision-Recall Curve for class {labelencoder.classes_[i]}')
387
388 plt.show()
389 y_val_pred_prob = nb.predict_proba(x_val)
390
391 for i in range(len(labelencoder.classes_)):
392     precision, recall, _ = precision_recall_curve(y_val == i, y_val_pred_prob[:, i])
393     PrecisionRecallDisplay(precision=precision, recall=recall, average_precision=np.mean(
394         precision)).plot()
395     plt.title(f'Precision-Recall Curve for class {labelencoder.classes_[i]}')
396
397 plt.show()
398
399 from sklearn.svm import SVC
400 from sklearn.metrics import PrecisionRecallDisplay
401
402 # svm = SVC(kernel='linear', probability=True)
403 # svm.fit(x_train, y_train)
404 from sklearn.model_selection import GridSearchCV
405 param_grid = {
406     'C': [0.1, 1, 10],
407     'kernel': ['linear', 'rbf'],
408     'gamma': [0.001, 0.01, 0.1]
409 }
410 grid_search = GridSearchCV(SVC(probability=True), param_grid, cv=5)
411 grid_search.fit(x_train, y_train)
412 best_model = grid_search.best_estimator_
413 # y_val_pred = svm.predict(x_val)
414 y_val_pred = best_model.predict(x_val)
415
416 val_accuracy = accuracy_score(y_val, y_val_pred)
417 print("Validation Accuracy:", val_accuracy)
418
419 val_cm = confusion_matrix(y_val, y_val_pred)
420 print("Validation Confusion Matrix:")
421 print(val_cm)
422
423 plot_confusion_matrix(val_cm, labels=labelencoder.classes_)
424
425 val_report = classification_report(y_val, y_val_pred)
426 print("Validation Classification Report:")
427 print(val_report)
428 y_val_pred_prob = best_model.predict_proba(x_val)
429 for i in range(len(labelencoder.classes_)):
430     precision, recall, _ = precision_recall_curve(y_val == i, y_val_pred_prob[:, i])
431     PrecisionRecallDisplay(precision=precision, recall=recall, average_precision=np.mean(
432         precision)).plot()

```

```

432 plt.title(f'Precision-Recall Curve for class {labelencoder.classes_[i]} (Validation)')
433
434 plt.show()
435
436 # y_test_pred = svm.predict(x_test)
437 y_test_pred = best_model.predict(x_test)
438 test_accuracy = accuracy_score(y_test, y_test_pred)
439 print("Test Accuracy:", test_accuracy)
440 test_cm = confusion_matrix(y_test, y_test_pred)
441 print("Test Confusion Matrix:")
442 print(test_cm)
443
444 plot_confusion_matrix(test_cm, labels=labelencoder.classes_)
445
446 test_report = classification_report(y_test, y_test_pred)
447 print("Test Classification Report:")
448 print(test_report)
449
450 y_test_pred_prob = best_model.predict_proba(x_test)
451 for i in range(len(labelencoder.classes_)):
452     precision, recall, _ = precision_recall_curve(y_test == i, y_test_pred_prob[:, i])
453     PrecisionRecallDisplay(precision=precision, recall=recall, average_precision=np.mean(
454         precision)).plot()
455     plt.title(f'Precision-Recall Curve for class {labelencoder.classes_[i]} (Test)')
456 plt.show()

```

C. Wage Dataset

```

1 from google.colab import drive
2 drive.mount('/content/drive')
3 import pandas as pd
4 import numpy as np
5 from sklearn.model_selection import train_test_split, GridSearchCV
6 from sklearn.preprocessing import OneHotEncoder
7 from sklearn.compose import ColumnTransformer
8 from sklearn.pipeline import Pipeline
9 from sklearn.metrics import mean_squared_error, r2_score
10 import matplotlib.pyplot as plt
11 df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/DT-Wage.csv')
12 from sklearn.preprocessing import StandardScaler
13
14 df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/DT-Wage.csv')
15
16 df.shape
17
18 df.info()
19 un = df['maritl'].unique()
20 print(un)
21
22 un = df['race'].unique()
23 print(un)
24
25 un = df['education'].unique()
26 print(un)
27
28 un = df['region'].unique()
29 print(un)
30
31 un = df['jobclass'].unique()
32 print(un)
33
34 un = df['health'].unique()

```

```

35 print(un)
36
37 categorical = df.select_dtypes(include=['object']).columns
38 numerical = df.select_dtypes(exclude=['object']).columns
39
40 encoder = OneHotEncoder(drop='first', sparse_output=False)
41 encoded_data = encoder.fit_transform(df[categorical])
42 encoded_df = pd.DataFrame(encoded_data, columns=encoder.get_feature_names_out(categorical))
43 encoded_df.index = df.index
44
45 df = df.drop(categorical, axis=1)
46 df = pd.concat([df, encoded_df], axis=1)
47
48 X = df.drop(columns=['wage'])
49 y = df['wage']
50 X.head()
51
52 X.shape
53
54 scaler = StandardScaler()
55 X_scaled = scaler.fit_transform(X)
56
57 X_train, X_temp, y_train, y_temp = train_test_split(X_scaled, y, test_size=0.3, random_state
    =42)
58 X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)
59
60 from sklearn.svm import SVR
61 from sklearn.model_selection import GridSearchCV
62 from sklearn.metrics import mean_squared_error, r2_score
63 svr = SVR()
64
65 param_grid = {
66     'C': [0.1, 1, 10, 100],
67     'kernel': ['linear', 'poly', 'rbf'],
68     'degree': [3, 4, 5],
69     'gamma': ['scale', 'auto'],
70     'epsilon': [0.1, 0.2, 0.5],
71 }
72
73 grid_search = GridSearchCV(svr, param_grid, cv=5, scoring='neg_mean_squared_error', n_jobs=-1)
74 grid_search.fit(X_train, y_train)
75 print("Best hyperparameters:", grid_search.best_params_)
76
77 best_svr = grid_search.best_estimator_
78 y_pred_val = best_svr.predict(X_val)
79 def custom_mse(y_true, y_pred):
80     return np.mean((y_true - y_pred) ** 2)
81
82 def custom_r2(y_true, y_pred):
83     ss_total = np.sum((y_true - np.mean(y_true)) ** 2)
84     ss_residual = np.sum((y_true - y_pred) ** 2)
85     return 1 - (ss_residual / ss_total)
86
87 mse = custom_mse(y_val, y_pred_val)
88 r2 = custom_r2(y_val, y_pred_val)
89 print(f"Validation MSE: {mse}")
90 print(f"Validation R^2: {r2}")
91
92
93 y_pred_test = best_svr.predict(X_test)
94 test_mse = custom_mse(y_test, y_pred_test)
95 test_r2 = custom_r2(y_test, y_pred_test)
96 print(f"Test MSE: {test_mse}")

```

```

97 print(f"Test R^2: {test_r2}")
98
99 plt.figure(figsize=(10, 6))
100 plt.scatter(y_val, y_pred_val, color='blue', alpha=0.5)
101 plt.plot([min(y_val), max(y_val)], [min(y_val), max(y_val)], color='red', linestyle='--')
102 plt.title("SVR: Actual vs Predicted (Validation Set)")
103 plt.xlabel("Actual Values")
104 plt.ylabel("Predicted Values")
105 plt.show()
106
107 residuals = y_val - y_pred_val
108 plt.figure(figsize=(10, 6))
109 plt.scatter(y_pred_val, residuals, color='green', alpha=0.5)
110 plt.hlines(0, xmin=min(y_pred_val), xmax=max(y_pred_val), colors='red', linestyle='--')
111 plt.title("Residuals Plot (Validation Set)")
112 plt.xlabel("Predicted Values")
113 plt.ylabel("Residuals")
114 plt.show()
115
116 plt.figure(figsize=(10, 6))
117 plt.scatter(y_test, y_pred_test, color='blue', alpha=0.5)
118 plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red', linestyle='--')
119 plt.title("SVR: Actual vs Predicted (Test Set)")
120 plt.xlabel("Actual Values")
121 plt.ylabel("Predicted Values")
122 plt.show()
123
124 residuals_test = y_test - y_pred_test
125 plt.figure(figsize=(10, 6))
126 plt.scatter(y_pred_test, residuals_test, color='green', alpha=0.5)
127 plt.hlines(0, xmin=min(y_pred_test), xmax=max(y_pred_test), colors='red', linestyle='--')
128 plt.title("Residuals Plot (Test Set)")
129 plt.xlabel("Predicted Values")
130 plt.ylabel("Residuals")
131 plt.show()

```

D. Credit Dataset

```

1 from google.colab import drive
2 drive.mount('/content/drive')
3 import pandas as pd
4 import numpy as np
5 from sklearn.model_selection import train_test_split, GridSearchCV
6 from sklearn.preprocessing import OneHotEncoder
7 from sklearn.compose import ColumnTransformer
8 from sklearn.pipeline import Pipeline
9 from sklearn.metrics import mean_squared_error, r2_score
10 import matplotlib.pyplot as plt
11 from sklearn.preprocessing import StandardScaler
12 df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/DT-Credit.csv')
13 df.head()
14 df.shape
15 df.describe()
16 df.isnull().sum()
17 un = df['Own'].unique()
18 print(un)
19
20 un = df['Student'].unique()
21 print(un)
22
23 un = df['Region'].unique()
24 print(un)
25

```



```

26 un = df['Married'].unique()
27 print(un)
28
29 categorical = df.select_dtypes(include=['object']).columns
30 numerical = df.select_dtypes(exclude=['object']).columns
31
32 encoder = OneHotEncoder(drop='first', sparse_output=False)
33 encoded_data = encoder.fit_transform(df[categorical])
34 encoded_df = pd.DataFrame(encoded_data, columns=encoder.get_feature_names_out(categorical))
35 encoded_df.index = df.index
36
37 df = df.drop(categorical, axis=1)
38 df = pd.concat([df, encoded_df], axis=1)
39 X = df.drop(columns=['Balance'])
40 y = df['Balance']
41 X.head()
42
43 X.shape
44
45 corr = df.corr()
46 target = corr['Balance'].sort_values(ascending=False)
47 print("Correlation of each feature with the target variable (Balance):")
48 print(target)
49
50 fig, ax = plt.subplots(1,1, figsize=(12, 12))
51 df.boxplot('Income', 'Balance', ax=ax)
52 plt.suptitle('Income vs Balance')
53 plt.title('')
54 plt.ylabel('Income')
55 plt.xticks(rotation=90)
56 plt.show()
57
58 fig, ax = plt.subplots(1,1, figsize=(12, 12))
59 df.boxplot('Age', 'Balance', ax=ax)
60 plt.suptitle('Age vs Balance')
61 plt.title('')
62 plt.ylabel('Age')
63 plt.xticks(rotation=90)
64 plt.show()
65
66 scaler = StandardScaler()
67 X_scaled = scaler.fit_transform(X)
68 X_train, X_temp, y_train, y_temp = train_test_split(X_scaled, y, test_size=0.3, random_state
    =42)
69 X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)
70
71 X_train.shape
72 X_val.shape
73 X_test.shape
74
75 from sklearn.model_selection import GridSearchCV
76 from sklearn.svm import SVR
77 from sklearn.metrics import mean_squared_error, r2_score
78
79 param_grid = {
80     'C': [0.1, 1, 10, 100,200,300,500,1500,2000],
81     'epsilon': [0.01, 0.1, 0.2,0.25,0.3],
82     'gamma': ['scale', 'auto']
83 }
84
85 svr = SVR(kernel='rbf')
86
87 grid_search = GridSearchCV(estimator=svr, param_grid=param_grid,

```

```

88         cv=5, scoring='neg_mean_squared_error', n_jobs=-1)
89
90 grid_search.fit(X_train, y_train)
91 best_params = grid_search.best_params_
92 print("Best hyperparameters found by GridSearchCV:")
93 print(best_params)
94
95 best_svr = grid_search.best_estimator_
96 y_pred_train = best_svr.predict(X_train)
97 y_pred_val = best_svr.predict(X_val)
98 y_pred_test = best_svr.predict(X_test)
99 def custom_mse(y_true, y_pred):
100     return np.mean((y_true - y_pred) ** 2)
101
102 def custom_r2(y_true, y_pred):
103     ss_total = np.sum((y_true - np.mean(y_true)) ** 2)
104     ss_residual = np.sum((y_true - y_pred) ** 2)
105     return 1 - (ss_residual / ss_total)
106
107 mse_train = custom_mse(y_train, y_pred_train)
108 mse_val = custom_mse(y_val, y_pred_val)
109 mse_test = custom_mse(y_test, y_pred_test)
110
111 r2_train = custom_r2(y_train, y_pred_train)
112 r2_val = custom_r2(y_val, y_pred_val)
113 r2_test = custom_r2(y_test, y_pred_test)
114
115 print("SVR Model with GridSearchCV Evaluation:")
116 print(f"Training MSE: {mse_train}")
117 print(f"Validation MSE: {mse_val}")
118 print(f"Test MSE: {mse_test}")
119 print(f"Training R : {r2_train}")
120 print(f"Validation R : {r2_val}")
121 print(f"Test R : {r2_test}")
122
123
124 plt.figure(figsize=(12, 4))
125 plt.subplot(1, 3, 1)
126 plt.scatter(y_train, y_pred_train, color='blue', alpha=0.5)
127 plt.plot([y_train.min(), y_train.max()], [y_train.min(), y_train.max()], color='red', lw=2)
128 plt.title('Training Set: Predicted vs Actual')
129 plt.xlabel('Actual')
130 plt.ylabel('Predicted')
131 plt.tight_layout()
132 plt.show()
133
134 plt.figure(figsize=(12, 4))
135 plt.subplot(1, 3, 2)
136 plt.scatter(y_val, y_pred_val, color='green', alpha=0.5)
137 plt.plot([y_val.min(), y_val.max()], [y_val.min(), y_val.max()], color='red', lw=2)
138 plt.title('Validation Set: Predicted vs Actual')
139 plt.xlabel('Actual')
140 plt.ylabel('Predicted')
141 plt.tight_layout()
142 plt.show()
143
144 plt.figure(figsize=(12, 4))
145 plt.subplot(1, 3, 3)
146 plt.scatter(y_test, y_pred_test, color='orange', alpha=0.5)
147 plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color='red', lw=2)
148 plt.title('Test Set: Predicted vs Actual')
149 plt.xlabel('Actual')
150 plt.ylabel('Predicted')

```

```
151
152 plt.tight_layout()
153 plt.show()
154 plt.figure(figsize=(12, 4))
155 plt.subplot(1, 3, 1)
156 plt.scatter(y_pred_train, y_pred_train - y_train, color='blue', alpha=0.5)
157 plt.hlines(y=0, xmin=y_pred_train.min(), xmax=y_pred_train.max(), color='red', lw=2)
158 plt.title('Training Set: Residuals')
159 plt.xlabel('Predicted')
160 plt.ylabel('Residuals (Error)')
161
162 plt.tight_layout()
163 plt.show()
164 plt.figure(figsize=(12, 4))
165 plt.subplot(1, 3, 2)
166 plt.scatter(y_pred_val, y_pred_val - y_val, color='green', alpha=0.5)
167 plt.hlines(y=0, xmin=y_pred_val.min(), xmax=y_pred_val.max(), color='red', lw=2)
168 plt.title('Validation Set: Residuals')
169 plt.xlabel('Predicted')
170 plt.ylabel('Residuals (Error)')
171
172 plt.tight_layout()
173 plt.show()
174 plt.figure(figsize=(12, 4))
175 plt.subplot(1, 3, 3)
176 plt.scatter(y_pred_test, y_pred_test - y_test, color='orange', alpha=0.5)
177 plt.hlines(y=0, xmin=y_pred_test.min(), xmax=y_pred_test.max(), color='red', lw=2)
178 plt.title('Test Set: Residuals')
179 plt.xlabel('Predicted')
180 plt.ylabel('Residuals (Error)')
181
182 plt.tight_layout()
183 plt.show()
```